

# GLYPH-X

a super regular RISC architecture

June 6, 2025 – DRAFT

**Michael Clark**

michaeljclark@mac.com



Maya glyphs in stucco on display at the Museo de Sitio in Palenque, Mexico.

*Image courtesy of Wikipedia. Public domain.*

to my mentors, whose guidance and wisdom made this work possible.

---

*"Perfection is achieved,  
not when there is nothing more to add,  
but when there is nothing left to take away."*

**— Antoine de Saint-Exupéry**

---

# Contents

|                                             |           |
|---------------------------------------------|-----------|
| <b>Preface</b>                              | <b>iv</b> |
| <b>1. Architecture</b>                      | <b>1</b>  |
| 1.1. Introduction . . . . .                 | 1         |
| 1.1.1. load-store . . . . .                 | 2         |
| 1.1.2. registers . . . . .                  | 2         |
| 1.1.3. memory . . . . .                     | 2         |
| 1.1.4. instructions . . . . .               | 3         |
| 1.1.5. constants . . . . .                  | 4         |
| 1.2. Core concepts . . . . .                | 5         |
| 1.3. Program streams . . . . .              | 6         |
| 1.4. Instruction format . . . . .           | 7         |
| 1.4.1. instruction templates . . . . .      | 7         |
| 1.4.2. instruction sizes . . . . .          | 8         |
| 1.4.3. instruction forms . . . . .          | 8         |
| 1.4.4. instruction decoding . . . . .       | 9         |
| 1.5. Register file . . . . .                | 10        |
| 1.6. Example pipeline . . . . .             | 11        |
| <b>2. Instructions</b>                      | <b>12</b> |
| 2.1. Instruction listing — 16-bit . . . . . | 12        |
| 2.1.1. break . . . . .                      | 12        |
| 2.1.2. j . . . . .                          | 12        |
| 2.1.3. b . . . . .                          | 12        |
| 2.1.4. ibj . . . . .                        | 13        |
| 2.1.5. link . . . . .                       | 13        |
| 2.1.6. movh . . . . .                       | 14        |
| 2.1.7. movw . . . . .                       | 14        |
| 2.1.8. movi . . . . .                       | 14        |
| 2.1.9. addi . . . . .                       | 14        |
| 2.1.10. srli . . . . .                      | 15        |
| 2.1.11. srai . . . . .                      | 15        |
| 2.1.12. slli . . . . .                      | 15        |
| 2.1.13. addh . . . . .                      | 15        |
| 2.1.14. leapc . . . . .                     | 16        |
| 2.1.15. loadpc . . . . .                    | 16        |
| 2.1.16. storepc . . . . .                   | 16        |

## Contents

|                                              |           |
|----------------------------------------------|-----------|
| 2.1.17. load . . . . .                       | 16        |
| 2.1.18. store . . . . .                      | 17        |
| 2.1.19. compare . . . . .                    | 17        |
| 2.1.20. logic . . . . .                      | 18        |
| 2.1.21. pin . . . . .                        | 18        |
| 2.1.22. and . . . . .                        | 19        |
| 2.1.23. or . . . . .                         | 19        |
| 2.1.24. xor . . . . .                        | 19        |
| 2.1.25. add . . . . .                        | 19        |
| 2.1.26. srl . . . . .                        | 20        |
| 2.1.27. sra . . . . .                        | 20        |
| 2.1.28. sll . . . . .                        | 20        |
| 2.1.29. sub . . . . .                        | 20        |
| 2.1.30. mul . . . . .                        | 21        |
| 2.1.31. div . . . . .                        | 21        |
| 2.1.32. illegal . . . . .                    | 21        |
| <b>3. Assembly</b>                           | <b>22</b> |
| 3.1. Introduction . . . . .                  | 22        |
| 3.2. Concepts . . . . .                      | 22        |
| 3.2.1. assembly . . . . .                    | 22        |
| 3.2.2. object . . . . .                      | 22        |
| 3.2.3. archive . . . . .                     | 22        |
| 3.2.4. executable . . . . .                  | 22        |
| 3.2.5. shared library . . . . .              | 23        |
| 3.2.6. section . . . . .                     | 23        |
| 3.2.7. segment . . . . .                     | 23        |
| 3.2.8. symbol . . . . .                      | 23        |
| 3.2.9. relocation . . . . .                  | 23        |
| 3.2.10. program linking . . . . .            | 23        |
| 3.3. Directives . . . . .                    | 24        |
| 3.4. Pseudo-instructions . . . . .           | 25        |
| 3.5. Calling convention . . . . .            | 26        |
| 3.5.1. calling convention — 16-bit . . . . . | 26        |
| <b>A. Appendix</b>                           | <b>27</b> |
| A.1. Opcode summary — 16-bit . . . . .       | 27        |
| <b>References</b>                            | <b>28</b> |

# Preface

this document introduces a RISC architecture designed with simplicity as its dominating principle, along with several features not yet present in mainstream RISC architectures:

- a compression technique for split instruction and constant streams.
- a variable-length instruction format designed for vectorized decoding.
- an efficient frontier between hardware synthesis and software translation.

the document begins with a general introduction to RISC architectures, including principal elements such as load-store operation, memory, registers, instructions, and constants. it then presents an overview of the proposed architecture, followed by sections on program streams, the variable-length instruction format, the instruction set, its assembly syntax, and the programming language calling convention.

this document does not simply present an architecture that is a reimplementa-tion of existing ideas; rather, it introduces a novel architecture with an instruction coding scheme optimized for efficient parallel decoding in hardware or software, combined with contemporary practice in instruction semantics.

this work is intended for computer architects, systems programmers, students of computing machinery, and anyone interested in the evolution of RISC design. by the end of the document, readers should be able to reason about computer programs at the instruction level and understand the logic behind the architecture’s design choices—with enough detail to implement the architecture in hardware or software.

# 1. Architecture

## 1.1. Introduction

this section gives a brief introduction to RISC architectures.

A RISC<sup>1</sup> machine is a type of general-purpose computer with the characteristic that it has a reduced set of instructions in contrast to a CISC<sup>2</sup> machine. A RISC machine is Turing complete meaning it can perform any computation that a Turing machine can, given enough time and memory. a Turing machine [12] is a theoretical model of computation.

a RISC machine has a set of instructions which comprise basic operations such as: *load-from-memory*, *store-to-memory*, *add*, *subtract*, *compare*, plus *conditional branch* and *unconditional branch* instructions et cetera; which one can imagine as a list of instructions on a paper tape. each instruction has an *opcode*, which is a unique binary pattern that identifies the *operation*, plus several *operands*, which are arguments to the instruction.

```
register-0 = load-from-memory at tape-address-0
register-1 = load-from-memory at tape-address-1
register-2 = add register-0 and register-1
            store-to-memory register-2 at tape-address-2
```

some instructions have operands that point to values inside of *registers* in a *register-file* which is like a close filing cabinet containing cards with numbers on them, and some of these numbers are addresses that point to values in *main-memory* which is like a larger but slower filing cabinet. some of these values are *immediate* values which are small numbers listed inside of the instructions on the paper tape.

the paper tape is just a way conceptualize a list of instructions stored in *main-memory*. there is a special register called *PC* short for *program counter*, which points to the current position on the tape. after each instruction executes the tape is advanced to the next instruction and the *program counter* is incremented, until it encounters a *branch* instruction which causes it to move forwards or backwards to a different position on the tape. branch instructions can be *conditional* or *unconditional*. conditional branches are selectively executed based on the results of a comparison instruction.

---

<sup>1</sup>Reduced Instruction Set Computer

<sup>2</sup>Complex Instruction Set Computer

## 1. Architecture

### 1.1.1. load-store

a *load-store* architecture [3, 4] is a way to characterize RISC architectures where most instructions have simple operands that point to values held in registers, plus *load* and *store* instructions to retrieve and commit values to main memory. a *load-store* architecture alleviates the need to add complex addressing modes, plus input-output to peripherals and secondary storage use *MMIO*<sup>3</sup> to avoid needing special *Input/Output* instructions.

### 1.1.2. registers

registers are temporary storage used to fill input and output operands for the *ALU*<sup>4</sup> before and after execution of instructions. registers are organized as a word-addressable store where each register number refers to  $XLEN \in \{64, 128\}$  bits of data. *XLEN* is a parameter that specifies the width of registers in bits.




Figure 1.1.: organization of register storage.

### 1.1.3. memory

main memory is primary storage which in modern computers is most likely *DRAM*<sup>5</sup>. main memory is organized as a byte-addressable store where each address refers to a byte which is 8-bits of data. *ALEN* is a parameter that specifies the width of addresses in bits.




Figure 1.2.: organization of main-memory storage.

<sup>3</sup>*MMIO* - Memory-mapped I/O.

<sup>4</sup>ALU – Arithmetic logic unit.

<sup>5</sup>*DRAM* – Dynamic random-access memory.



## 1. Architecture

### 1.1.4. instructions

instruction memory is a memory region dedicated to program instructions. it is organized as a word-addressable store, similar to main memory, but is read-only and reserved exclusively for instructions. this configuration is typically referred to as a Harvard architecture [1], in contrast to a Von Neumann architecture [8, 9], which uses a shared memory region for both program instructions and data.

instruction words contain packets of 16-bits which are composed of size, opcode, and operand fields. instruction packets can be appended together to form larger instruction words with larger opcode and operand fields. section 1.4 provides details on the instruction template and instruction forms.

|                   |   |   |                  |   |              |
|-------------------|---|---|------------------|---|--------------|
| 15                | 7 | 6 | 2                | 1 | 0            |
| $operand_{[8:0]}$ |   |   | $opcode_{[4:0]}$ |   | $sz_{[1:0]}$ |

Figure 1.3.: 16-bit instruction containing one packet.

|                    |   |   |                  |   |              |
|--------------------|---|---|------------------|---|--------------|
| 15                 | 7 | 6 | 2                | 1 | 0            |
| $operand_{[8:0]}$  |   |   | $opcode_{[4:0]}$ |   | $sz_{[1:0]}$ |
| $operand_{[17:9]}$ |   |   | $opcode_{[9:5]}$ |   | $sz_{[3:2]}$ |

Figure 1.4.: 32-bit instruction containing two packets.

instruction words are packed together into instruction blocks containing instruction size information that is used to fuse them together to compose variable-length instructions [7]. instruction words can begin on any 16-bit aligned address. instruction memory is addressed with displacements from the *program counter* called *pc-relative* addresses.




Figure 1.5.: organization of instruction-memory with 16-bit instructions.




Figure 1.6.: organization of instruction-memory with 16-bit and 32-bit instructions.

## 1. Architecture

### 1.1.5. constants

constant memory [11] is a memory region dedicated to constants. it is organized as a word-addressable store like main-memory, but is read-only and restricted to constants. instructions can refer to constants using *ib-relative* addresses inside instruction operands. section 1.3 provides more details on the split instruction and constant streams.

constant memory is addressed with displacements from the *immediate base* register called *ib-relative* addresses which are stored inside instruction operand slots, are sized, scaled and aligned addresses computed relative to the *immediate base* register, which is like a program counter for constants. these diagrams show  $ib8(n)$  through  $ib64(n)$  *ib-relative* addresses for access to 8-bit through to 64-bit constants respectively: *note: ib-relative addresses alias a single constant storage space containing all types.*




Figure 1.7.: organization of *ib8* constant-memory storage.




Figure 1.8.: organization of *ib16* constant-memory storage.




Figure 1.9.: organization of *ib32* constant-memory storage.




Figure 1.10.: organization of *ib64* constant-memory storage.

### 1.2. Core concepts

glyph is a super regular RISC architecture that encodes constants in a secondary stream accessed via an *immediate base* register that points at immediate blocks containing constants accessed via a constant address mode. the *immediate base* register branches like the *program counter*, and procedure calls and returns set and restore (*pc,ib*) together.

glyph uses relative address vectors in its link register which is different to typical RISC architectures. glyph does this so that the branch instructions can fit (*pc,ib*) into the a single link register for compatibility with traditional RISC architectures. glyph achieves this by packing two relative (*pc,ib*) displacements into a relative address vector<sup>6</sup>.

immediate blocks can be switched using the immediate block branch instruction. immediate blocks, unlike typical RISC architectures, mean that most relocations are word sized like CISC architectures, and can use C-style structure packing and alignment rules.

this list outlines some differentiating elements of the super regular RISC architecture:

- variable length instruction format supporting 16, 32, 64, and 128-bit instructions.
- 16-bit compressed instruction packets can access 8 registers.
- (*pc,ib*) is a *program counter* and *immediate base* register address vector.
- link register contains a packed relative (*pc,ib*) address vector to function entry.
- *ibj immediate-block-jump* adds a relative address to the *immediate base* register.
- *movw move-word-immediate-block* uses an unsigned displacement to access constants.
- *jalib jump-and-link-immediate-block* or *call link function* links address vector and adds constants to (*pc,ib*) and is used to branch the *program counter* and *immediate base* register at the same time for calling procedures.
- *jtlb jump-to-link-immediate-block* or *return link function* subtracts link vector from and adds constants to (*pc,ib*) and is used to branch the *program counter* and *immediate base* register at the same time for returning from called procedures.
- *pin pack-indirect* packs two absolute addresses as relative address vector from (*pc,ib*) and is used for calling absolute addresses such as virtual functions.

the *pin* and *link* instructions form a modular arithmetic ring of relative address vectors, due to their use of relative addresses. relative address vectors can be encrypted, decrypted, and authenticated to form a verifiable chain of addresses for control flow integrity.

---

<sup>6</sup>the architecture defines two parameters: *ALEN* and *XLEN*, which respectively refer to width of addresses and width of general purpose registers in bits. when  $XLEN \geq ALEN \times 2$  it is possible to pack absolute addresses instead of relative addresses, as would be the case where  $ALEN=64$  and  $XLEN=128$ .

### 1.3. Program streams

glyph separates the stored programs into two streams, one with instructions and one with constants [11]. the instruction stream is addressed with the *program counter* (*pc*) and the constant stream is addressed with the *immediate base* (*ib*) register. this could be referred to as split program stream Harvard architecture.




Figure 1.11.: program counter and immediate base register.

instruction blocks are aligned memory blocks addressed by the *program counter*. instruction memory is addressed with *pc-relative* addresses in instruction immediate values or indirectly with constants accessed using *ib-relative* addresses stored inside instruction operand slots.

immediate blocks are aligned memory blocks addressed by the *immediate base* register. constant memory is addressed with *ib-relative* addresses stored inside instruction operand slots, which are sized, scaled, and aligned relative to the *immediate base* register.

the constant stream branches independently using the *immediate-block-jump* instruction to update the *immediate base* register, or together with the *program counter* in procedure call and return instructions that branch instruction and constants at the same time; *jump-and-link-immediate-block* and *jump-to-link-immediate-block* add and subtract address vectors to (*pc,ib*). the *pack-indirect* instruction allows absolute (*pc,ib*) addresses to be packed into an address vector for indirect calls.

the instruction forms use bonded operand slots for immediate operands and operand bits do not overlap opcode bits. the use of immediate blocks means large immediate constants can all be accessed using short references encoded inside of operand slots for instructions that use an immediate block relative addressing mode.

## 1.4. Instruction format

instructions use a variable length format [7] with a 16-bit instruction packet that supports 16, 32, 64, and 128-bit instruction words. the instruction packet uses a *super regular* scheme, whereby successive instruction packets extend the fields in the previous instruction packet [6]. instruction words can begin on any 16-bit aligned address.

### 1.4.1. instruction templates

the variable length instruction format has a single base format where fields in the template instruction form are extended by successive instruction packets.

|                                 |   |                                |   |                            |   |
|---------------------------------|---|--------------------------------|---|----------------------------|---|
| 15                              | 7 | 6                              | 2 | 1                          | 0 |
| <i>operand</i> <sub>[8:0]</sub> |   | <i>opcode</i> <sub>[4:0]</sub> |   | <i>sz</i> <sub>[1:0]</sub> |   |

Figure 1.12.: instruction template — 16-bit.

|                                  |   |                                |   |                            |   |
|----------------------------------|---|--------------------------------|---|----------------------------|---|
| 15                               | 7 | 6                              | 2 | 1                          | 0 |
| <i>operand</i> <sub>[8:0]</sub>  |   | <i>opcode</i> <sub>[4:0]</sub> |   | <i>sz</i> <sub>[1:0]</sub> |   |
| <i>operand</i> <sub>[17:9]</sub> |   | <i>opcode</i> <sub>[9:5]</sub> |   | <i>sz</i> <sub>[3:2]</sub> |   |

Figure 1.13.: instruction template — 32-bit.

|                                   |   |                                  |   |                            |   |
|-----------------------------------|---|----------------------------------|---|----------------------------|---|
| 15                                | 7 | 6                                | 2 | 1                          | 0 |
| <i>operand</i> <sub>[8:0]</sub>   |   | <i>opcode</i> <sub>[4:0]</sub>   |   | <i>sz</i> <sub>[1:0]</sub> |   |
| <i>operand</i> <sub>[17:9]</sub>  |   | <i>opcode</i> <sub>[9:5]</sub>   |   | <i>sz</i> <sub>[3:2]</sub> |   |
| <i>operand</i> <sub>[26:18]</sub> |   | <i>opcode</i> <sub>[14:10]</sub> |   | <i>sz</i> <sub>[5:4]</sub> |   |
| <i>operand</i> <sub>[35:27]</sub> |   | <i>opcode</i> <sub>[19:15]</sub> |   | <i>sz</i> <sub>[7:6]</sub> |   |

Figure 1.14.: instruction template — 64-bit.

|                                   |   |                                  |   |                              |   |
|-----------------------------------|---|----------------------------------|---|------------------------------|---|
| 15                                | 7 | 6                                | 2 | 1                            | 0 |
| <i>operand</i> <sub>[8:0]</sub>   |   | <i>opcode</i> <sub>[4:0]</sub>   |   | <i>sz</i> <sub>[1:0]</sub>   |   |
| <i>operand</i> <sub>[17:9]</sub>  |   | <i>opcode</i> <sub>[9:5]</sub>   |   | <i>sz</i> <sub>[3:2]</sub>   |   |
| <i>operand</i> <sub>[26:18]</sub> |   | <i>opcode</i> <sub>[14:10]</sub> |   | <i>sz</i> <sub>[5:4]</sub>   |   |
| <i>operand</i> <sub>[35:27]</sub> |   | <i>opcode</i> <sub>[19:15]</sub> |   | <i>sz</i> <sub>[7:6]</sub>   |   |
| <i>operand</i> <sub>[44:36]</sub> |   | <i>opcode</i> <sub>[24:20]</sub> |   | <i>sz</i> <sub>[9:8]</sub>   |   |
| <i>operand</i> <sub>[53:45]</sub> |   | <i>opcode</i> <sub>[29:25]</sub> |   | <i>sz</i> <sub>[11:10]</sub> |   |
| <i>operand</i> <sub>[62:54]</sub> |   | <i>opcode</i> <sub>[34:30]</sub> |   | <i>sz</i> <sub>[13:12]</sub> |   |
| <i>operand</i> <sub>[71:63]</sub> |   | <i>opcode</i> <sub>[39:35]</sub> |   | <i>sz</i> <sub>[15:14]</sub> |   |

Figure 1.15.: instruction template — 128-bit.

## 1. Architecture

### 1.4.2. instruction sizes

the variable length instruction format has a 2-bit size field in a fixed position in every 16-bit instruction packet, to reduce the complexity of size-decoding for variable length instructions. this table shows the size fields for the different instruction sizes:

| Instruction size | Size vector                      |
|------------------|----------------------------------|
| 16-bit           | {00}                             |
| 32-bit           | {01, 11}                         |
| 64-bit           | {10, 11, 11, 11}                 |
| 128-bit          | {11, 11, 11, 11, 11, 11, 11, 11} |

Table 1.1.: variable-length instruction size fields

### 1.4.3. instruction forms

the instructions forms are super regular in that operand and opcode bits do not overlap and the number and complexity of the formats is reduced so that vectorized instruction decoding is simple in both hardware and software. the scheme is designed so that 1-bit of coding space in the larger packet can be used to extend register sizes.

#### 16-bit instruction forms

this section details the layout of the 16-bit instruction forms:

|                             |   |   |   |                                |   |
|-----------------------------|---|---|---|--------------------------------|---|
| 15                          | 7 | 6 | 2 | 1                              | 0 |
| <i>imm</i> <sub>[5:0]</sub> |   |   |   | <i>opcode</i> <sub>[4:0]</sub> |   |
|                             |   |   |   | <b>00</b>                      |   |

Figure 1.16.: 16-bit large immediate.

|                            |    |                             |   |   |                                |   |           |
|----------------------------|----|-----------------------------|---|---|--------------------------------|---|-----------|
| 15                         | 13 | 12                          | 7 | 6 | 2                              | 1 | 0         |
| <i>rc</i> <sub>[2:0]</sub> |    | <i>imm</i> <sub>[5:0]</sub> |   |   | <i>opcode</i> <sub>[4:0]</sub> |   | <b>00</b> |

Figure 1.17.: 16-bit one operand with immediate.

|                            |    |                            |    |                             |   |                                |   |           |   |
|----------------------------|----|----------------------------|----|-----------------------------|---|--------------------------------|---|-----------|---|
| 15                         | 13 | 12                         | 10 | 9                           | 7 | 6                              | 2 | 1         | 0 |
| <i>rc</i> <sub>[2:0]</sub> |    | <i>rb</i> <sub>[2:0]</sub> |    | <i>imm</i> <sub>[2:0]</sub> |   | <i>opcode</i> <sub>[4:0]</sub> |   | <b>00</b> |   |

Figure 1.18.: 16-bit two operand with immediate.

|                            |    |                            |    |                            |   |                                |   |           |   |
|----------------------------|----|----------------------------|----|----------------------------|---|--------------------------------|---|-----------|---|
| 15                         | 13 | 12                         | 10 | 9                          | 7 | 6                              | 2 | 1         | 0 |
| <i>rc</i> <sub>[2:0]</sub> |    | <i>rb</i> <sub>[2:0]</sub> |    | <i>ra</i> <sub>[2:0]</sub> |   | <i>opcode</i> <sub>[4:0]</sub> |   | <b>00</b> |   |

Figure 1.19.: 16-bit three operand.

## 1. Architecture

### 32-bit instruction forms

this section details the layout of the 32-bit instruction forms:

|                              |   |                                |   |           |   |
|------------------------------|---|--------------------------------|---|-----------|---|
| 15                           | 7 | 6                              | 2 | 1         | 0 |
| <i>imm</i> <sub>[8:0]</sub>  |   | <i>opcode</i> <sub>[4:0]</sub> |   | <b>01</b> |   |
| <i>imm</i> <sub>[17:9]</sub> |   | <i>opcode</i> <sub>[9:5]</sub> |   | <b>11</b> |   |

Figure 1.20.: 32-bit large immediate.

|                            |    |                              |   |                                |   |           |   |
|----------------------------|----|------------------------------|---|--------------------------------|---|-----------|---|
| 15                         | 13 | 12                           | 7 | 6                              | 2 | 1         | 0 |
| <i>rc</i> <sub>[2:0]</sub> |    | <i>imm</i> <sub>[5:0]</sub>  |   | <i>opcode</i> <sub>[4:0]</sub> |   | <b>01</b> |   |
| <i>rc</i> <sub>[5:3]</sub> |    | <i>imm</i> <sub>[11:6]</sub> |   | <i>opcode</i> <sub>[9:5]</sub> |   | <b>11</b> |   |

Figure 1.21.: 32-bit one operand with immediate.

|                            |    |                            |    |                             |   |                                |   |           |   |
|----------------------------|----|----------------------------|----|-----------------------------|---|--------------------------------|---|-----------|---|
| 15                         | 13 | 12                         | 10 | 9                           | 7 | 6                              | 2 | 1         | 0 |
| <i>rc</i> <sub>[2:0]</sub> |    | <i>rb</i> <sub>[2:0]</sub> |    | <i>imm</i> <sub>[2:0]</sub> |   | <i>opcode</i> <sub>[4:0]</sub> |   | <b>01</b> |   |
| <i>rc</i> <sub>[5:3]</sub> |    | <i>rb</i> <sub>[5:3]</sub> |    | <i>imm</i> <sub>[5:3]</sub> |   | <i>opcode</i> <sub>[9:5]</sub> |   | <b>11</b> |   |

Figure 1.22.: 32-bit two operand with immediate.

|                            |    |                            |    |                            |   |                                |   |           |   |
|----------------------------|----|----------------------------|----|----------------------------|---|--------------------------------|---|-----------|---|
| 15                         | 13 | 12                         | 10 | 9                          | 7 | 6                              | 2 | 1         | 0 |
| <i>rc</i> <sub>[2:0]</sub> |    | <i>rb</i> <sub>[2:0]</sub> |    | <i>ra</i> <sub>[2:0]</sub> |   | <i>opcode</i> <sub>[4:0]</sub> |   | <b>01</b> |   |
| <i>rc</i> <sub>[5:3]</sub> |    | <i>rb</i> <sub>[5:3]</sub> |    | <i>ra</i> <sub>[5:3]</sub> |   | <i>opcode</i> <sub>[9:5]</sub> |   | <b>11</b> |   |

Figure 1.23.: 32-bit three operand.

#### 1.4.4. instruction decoding

the instruction size encoding and field extension scheme has been designed to minimize complexity for parallel decoding in hardware and software. the following table shows the instruction width combinations for various width parallel instruction decoders.

| Packets | Bits    | Bytes | $n_{combos}$ | $\log_2(n)$ |
|---------|---------|-------|--------------|-------------|
| 1-wide  | 16-bit  | 2     | 1            | 0           |
| 2-wide  | 32-bit  | 4     | 16           | 4           |
| 4-wide  | 64-bit  | 8     | 58           | 5.86        |
| 8-wide  | 128-bit | 16    | 574          | 9.16        |
| 16-wide | 256-bit | 32    | 51904        | 15.66       |

Table 1.2.: instruction decode combinations

## 1.5. Register file

the glyph register file is extensible due to the variable length instruction format and supports a different number of registers depending on the instruction size.

- 16-bit instruction packet can access 8 registers with up to 3 operands.
- 32-bit instruction packet can access 64 registers with up to 3 operands.
- 64-bit instruction packet can access 64 registers with up to 6 operands.

the register state accessible by the 16-bit instruction packet is comprised of:

- *program counter* register (aligned to 2 bytes).
- *immediate base* register (aligned to 64 bytes).
- 8 × general purpose registers (*r0* through *r7*).
- 1 × predicate register (*flag*).

the following diagram shows the register state accessible by the 16-bit instruction packet:




Figure 1.24.: register state accessible by 16-bit instruction packet.

the use of  $ALEN^7$  and  $XLEN^8$  parameters is to indicate that the width of addresses can be less than the width of the general purpose registers.

<sup>7</sup> $ALEN$  refers to the width of addresses in bits.

<sup>8</sup> $XLEN$  refers to the width of general purpose registers in bits.



## 1.6. Example pipeline

an illustrative micro-architecture is proposed based on the classic 5-stage RISC micro-architecture [10] with the addition of an *operand fetch* stage and a *constant memory* port. this revised 6-stage micro-architecture is composed of the following pipeline stages:

- IF — *instruction fetch*: reads instructions from memory into a fetch buffer.
- ID — *instruction decode*: decodes instruction length, opcode, and operands.
- OF — *operand fetch*: reads operands from register file and constant memory.
- EX — *execute*: performs logical operations or arithmetic on the operands.
- MA — *memory access*: loads data from or stores data to memory.
- WB — *writeback*: writes results back to the register file.

a simplified micro-architecture using those pipeline stages might look like this: this example omits hazard detection and forwarding logic for the sake of simplicity.




Figure 1.25.: sample 6-stage micro-architecture with support for constant memory.

## 2. Instructions

### 2.1. Instruction listing — 16-bit

#### 2.1.1. break

|             |   |               |             |   |   |
|-------------|---|---------------|-------------|---|---|
| 15          | 7 | 6             | 2           | 1 | 0 |
| <i>imm9</i> |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>uimm</b> |   | <b>00000</b>  | <b>00</b>   |   |   |

**break** *uimm9*

the *break* instruction causes a debugger trap. *program counter* and trap cause are saved to privileged registers for the operating system to dispatch to a debugger and the *program counter* is set to a trap vector address.

#### 2.1.2. j

|             |   |               |             |   |   |
|-------------|---|---------------|-------------|---|---|
| 15          | 7 | 6             | 2           | 1 | 0 |
| <i>imm9</i> |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>simm</b> |   | <b>00001</b>  | <b>00</b>   |   |   |

**j** *simm9* × 2

the *j* or *jump* instruction is an unconditional branch instruction that adds a relative immediate address to the *program counter*. the resulting *program counter* address is [*pc* + *simm9* × 2].

$$pc = pc + simm9 * 2$$

#### 2.1.3. b

|             |   |               |             |   |   |
|-------------|---|---------------|-------------|---|---|
| 15          | 7 | 6             | 2           | 1 | 0 |
| <i>imm9</i> |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>simm</b> |   | <b>00010</b>  | <b>00</b>   |   |   |

**b** *simm9* × 2

the *b* or *branch* instruction is a conditional branch instruction that adds a relative immediate address to the *program counter*. if the *flag* register has been set by a compare instruction, the resulting *program counter* address is [*pc* + *simm9* × 2], otherwise the *program counter* is advanced normally.

```
if flag:
    pc = pc + simm9 * 2
```

## 2. Instructions

### 2.1.4. ibj

|             |   |   |               |             |   |
|-------------|---|---|---------------|-------------|---|
| 15          | 7 | 6 | 2             | 1           | 0 |
| <i>imm9</i> |   |   | <i>opcode</i> | <i>size</i> |   |
| <b>simm</b> |   |   | <b>00011</b>  | <b>00</b>   |   |

**ibj**  $\text{simm9} \times 64$

the *ibj* or *immediate-block-jump* instruction adds a 64-bit relative address to the *immediate base* register. the resulting *immediate base* address is  $[ib + \text{simm9} \times 64]$ .

$$ib = ib + \text{simm9} * 64$$

### 2.1.5. link

|            |    |             |   |               |             |   |   |
|------------|----|-------------|---|---------------|-------------|---|---|
| 15         | 13 | 12          | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i>  |    | <i>imm6</i> |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>fun</b> |    | <b>uimm</b> |   | <b>00100</b>  | <b>00</b>   |   |   |

**link.i64**  $\text{fun3}, \text{ib64}(\text{uimm6})$

the *link* instruction loads a 64-bit constant addressed by  $[ib + \text{uimm6} \times 8]$  containing a *i32x2* relative address vector, which it adds it to  $(pc, ib)$ , conditionally subtracts the source link register (*r6* or *r7*), then conditionally adds the difference to the destination link register (*r6* or *r7*), depending on the value of *fun3*. the type of linkage in *fun3* can be one of:

| value | mnemonic | description       | link-type | dst-link | src-link |
|-------|----------|-------------------|-----------|----------|----------|
| 0     | jib      | jump              | jump      | -        | -        |
| 1     | -        | -                 | -         | -        | -        |
| 2     | jalib    | jump-and-link     | call      | r6       | -        |
| 3     | jalib    | jump-and-link     | call      | r7       | -        |
| 4     | jtlib    | jump-to-link      | ret       | -        | r6       |
| 5     | jtlib    | jump-to-link      | ret       | -        | r7       |
| 6     | jalaib   | jump-and-link-add | tail      | r6       | -        |
| 7     | jalaib   | jump-and-link-add | tail      | r7       | -        |

Table 2.1.: Link instruction functions

```

lr = 0b110 + (fun3 & 1)
cval = const-mem<i32x2>[ib + uimm6 * 8]
dval = fun3 == jalaib ? cval + auth-decrypt(reg<i32x2>[lr]) : cval
lval = fun3 == jtlib ?      auth-decrypt(reg<i32x2>[lr]) : 0
(dpc,dib) = dval
(lpc,lib) = lval
pc = pc + dpc - lpc
ib = ib + dib - lib
if fun3 == jalib or fun3 == jalaib:
    reg[lr] = auth-encrypt(reg<i32x2>(dpc,dib))

```

the *auth-encrypt* and *auth-decrypt* functions are placeholders for authenticated encryption, and decryption of relative address vectors. authentication failures should generate a trap for the operating system to dispatch to a control flow integrity trap handler.

## 2. Instructions

### 2.1.6. movh

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>uimm</b> |    |   | <b>00101</b>  | <b>00</b>   |   |   |

**movh.i64** *rc*, *ib32*(*uimm6*)

the *movh* or *move-half-immediate-block* instruction loads a 32-bit constant addressed by  $[ib + uimm6 \times 4]$ , which it sign-extends to 64-bits, then saves to the *rc* register.

$$\text{reg}[rc] = \text{const-mem}<i32>[ib + uimm6 * 4]$$

### 2.1.7. movw

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>uimm</b> |    |   | <b>00110</b>  | <b>00</b>   |   |   |

**movw.i64** *rc*, *ib64*(*uimm6*)

the *movw* or *move-word-immediate-block* instruction loads a 64-bit constant addressed by  $[ib + uimm6 \times 8]$  then saves it to the *rc* register.

$$\text{reg}[rc] = \text{const-mem}<i64>[ib + uimm6 * 8]$$

### 2.1.8. movi

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>simm</b> |    |   | <b>00111</b>  | <b>00</b>   |   |   |

**movi.i64** *rc*, *simm6*

the *movi* or *move-immediate* instruction sign-extends the immediate value in *simm6* then saves the result in the *rc* register.

$$\text{reg}[rc] = \text{simm6}$$

### 2.1.9. addi

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>simm</b> |    |   | <b>01000</b>  | <b>00</b>   |   |   |

**addi.i64** *rc*, *simm6*

the *addi* or *add-immediate* instruction sign-extends the immediate value in *simm6* then adds it to the *rc* register.

$$\text{reg}[rc] = \text{reg}[rc] + \text{simm6}$$

## 2. Instructions

### 2.1.10. srli

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>uimm</b> |    |   | <b>01001</b>  | <b>00</b>   |   |   |

**srli.i64** *rc*, *uimm6*

the *srli* or *shift-right-logical-immediate* instruction performs a logical right shift by *uimm6* bits of the value in the *rc* register. zeros are copied into the left most bits.

$$\text{reg}[\text{rc}] = \text{reg}\langle\text{u64}\rangle[\text{rc}] \gg \text{uimm6}$$

### 2.1.11. srai

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>uimm</b> |    |   | <b>01010</b>  | <b>00</b>   |   |   |

**srai.i64** *rc*, *uimm6*

the *srai* or *shift-right-arithmetic-immediate* instruction performs an arithmetic right shift by *uimm6* bits of the value in the *rc* register. the sign is copied into the left most bits.

$$\text{reg}[\text{rc}] = \text{reg}\langle\text{i64}\rangle[\text{rc}] \gg \text{uimm6}$$

### 2.1.12. slli

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>uimm</b> |    |   | <b>01011</b>  | <b>00</b>   |   |   |

**slli.i64** *rc*, *uimm6*

the *slli* or *shift-left-logical-immediate* instruction performs a logical left shift by *uimm6* bits of the value in the *rc* register. zeros are copied into the right most bits.

$$\text{reg}[\text{rc}] = \text{reg}[\text{rc}] \ll \text{uimm6}$$

### 2.1.13. addh

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>uimm</b> |    |   | <b>01100</b>  | <b>00</b>   |   |   |

**addh.i64** *rc*, *ib32*(*uimm6*)

the *addh* or *add-half-immediate-block* instruction loads a 32-bit constant addressed by [*ib* + *uimm6* × 4], which it sign-extends to 64-bits, then adds to the *rc* register.

$$\text{reg}[\text{rc}] = \text{reg}[\text{rc}] + \text{const-mem}\langle\text{i32}\rangle[\text{ib} + \text{uimm6} * 4]$$

## 2. Instructions

### 2.1.14. leapc

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>uimm</b> |    |   | <b>01101</b>  | <b>00</b>   |   |   |

**leapc.i64** *rc*, *ib*32(*uimm6*)(*pc*)

the *leapc* or *load-effective-address-pc* instruction loads a 32-bit constant addressed by [*ib* + *uimm6* × 4], which it sign-extends to 64-bits, adds it to the *program counter*, and saves the result in the *rc* register.

$$\text{reg}[\text{rc}] = \text{pc} + \text{const-mem}\langle\text{i32}\rangle[\text{ib} + \text{uimm6} * 4]$$

### 2.1.15. loadpc

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>uimm</b> |    |   | <b>01110</b>  | <b>00</b>   |   |   |

**loadpc.i64** *rc*, *ib*32(*uimm6*)(*pc*)

the *loadpc* instruction loads a 32-bit constant addressed by [*ib* + *uimm6* × 4] which it sign-extends to 64-bit, adds it to the *program counter* to form an address, then loads a 64-bit value from memory at that address and saves the result in the *rc* register.

$$\text{reg}[\text{rc}] = \text{mem}\langle\text{i64}\rangle[\text{pc} + \text{const-mem}\langle\text{i32}\rangle[\text{ib} + \text{uimm6} * 4]]$$

### 2.1.16. storepc

|           |             |    |   |               |             |   |   |
|-----------|-------------|----|---|---------------|-------------|---|---|
| 15        | 13          | 12 | 7 | 6             | 2           | 1 | 0 |
| <i>rc</i> | <i>imm6</i> |    |   | <i>opcode</i> | <i>size</i> |   |   |
| <b>rc</b> | <b>uimm</b> |    |   | <b>01111</b>  | <b>00</b>   |   |   |

**storepc.i64** *rc*, *ib*32(*uimm6*)(*pc*)

the *storepc* instruction loads a 32-bit constant addressed by [*ib* + *uimm6* × 4] which it sign-extends to 64-bit, adds it to the *program counter* to form an address, then stores to memory at that address a 64-bit value from the *rc* register.

$$\text{mem}\langle\text{i64}\rangle[\text{pc} + \text{const-mem}\langle\text{i32}\rangle[\text{ib} + \text{uimm6} * 4]] = \text{reg}[\text{rc}]$$

### 2.1.17. load

|           |           |             |    |               |   |             |   |   |   |
|-----------|-----------|-------------|----|---------------|---|-------------|---|---|---|
| 15        | 13        | 12          | 10 | 9             | 7 | 6           | 2 | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>imm3</i> |    | <i>opcode</i> |   | <i>size</i> |   |   |   |
| <b>rc</b> | <b>rb</b> | <b>uimm</b> |    | <b>10000</b>  |   | <b>00</b>   |   |   |   |

**load.i64** *rc*, (*uimm3* × 8)(*rb*)

the *load* instruction computes the address [*rb* + *uimm3* × 8] then loads a 64-bit value from memory at that address and saves the result in the *rc* register.

$$\text{reg}[\text{rc}] = \text{mem}\langle\text{i64}\rangle[\text{reg}[\text{rb}] + \text{uimm3} * 8]$$

## 2. Instructions

### 2.1.18. store

|           |           |             |               |             |   |   |   |   |   |
|-----------|-----------|-------------|---------------|-------------|---|---|---|---|---|
| 15        | 13        | 12          | 10            | 9           | 7 | 6 | 2 | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>imm3</i> | <i>opcode</i> | <i>size</i> |   |   |   |   |   |
| <b>rc</b> | <b>rb</b> | <b>uimm</b> | <b>10001</b>  | <b>00</b>   |   |   |   |   |   |

**store.i64** *rc*, (*uimm3* × 8)(*rb*)

the *store* instruction computes the address [*rb* + *uimm3* × 8] then stores a 64-bit value to memory at that address containing a 64-bit value from the *rc* register.

$$\text{mem}\langle\text{i64}\rangle[\text{reg}[\text{rb}] + \text{uimm3} * 8] = \text{reg}[\text{rc}]$$

### 2.1.19. compare

|           |           |             |               |             |   |   |   |   |   |
|-----------|-----------|-------------|---------------|-------------|---|---|---|---|---|
| 15        | 13        | 12          | 10            | 9           | 7 | 6 | 2 | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>imm3</i> | <i>opcode</i> | <i>size</i> |   |   |   |   |   |
| <b>rc</b> | <b>rb</b> | <b>fun</b>  | <b>10010</b>  | <b>00</b>   |   |   |   |   |   |

**cmp.i64** *rc*, *rb*, *fun3*

the *compare* instruction performs a comparison between the value in *rb* and *rc* then saves the result in the *flag* register. the *compare* opcode is also used to perform conditional move whereby the *rb* register is copied into the *rc* if the *flag* register is set. the type of comparison in *fun3* can be one of:

| value | mnemonic     | description                        |
|-------|--------------|------------------------------------|
| 0     | <i>lt</i>    | <i>less than (signed)</i>          |
| 1     | <i>ge</i>    | <i>greater or equal (signed)</i>   |
| 2     | <i>eq</i>    | <i>equal</i>                       |
| 3     | <i>ne</i>    | <i>not equal</i>                   |
| 4     | <i>ltu</i>   | <i>less than (unsigned)</i>        |
| 5     | <i>geu</i>   | <i>greater or equal (unsigned)</i> |
| 6     | <i>cmov</i>  | <i>conditional move</i>            |
| 7     | <i>ncmov</i> | <i>negated conditional move</i>    |

Table 2.2.: Compare instruction functions

```
match fun3
| lt    -> flag = reg<i64>[rc] < reg<i64>[rb]
| ge    -> flag = reg<i64>[rc] >= reg<i64>[rb]
| eq    -> flag = reg[rc] == reg[rb]
| ne    -> flag = reg[rc] != reg[rb]
| ltu   -> flag = reg<u64>[rc] < reg<u64>[rb]
| geu   -> flag = reg<u64>[rc] >= reg<u64>[rb]
| cmov  -> if flag:
            reg[rc] = reg[rb]
| ncmov -> if not flag:
            reg[rc] = reg[rb]
```

## 2. Instructions

### 2.1.20. logic

|           |           |             |               |             |                             |   |   |   |   |
|-----------|-----------|-------------|---------------|-------------|-----------------------------|---|---|---|---|
| 15        | 13        | 12          | 10            | 9           | 7                           | 6 | 2 | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>imm3</i> | <i>opcode</i> | <i>size</i> |                             |   |   |   |   |
| <b>rc</b> | <b>rb</b> | <b>fun</b>  | <b>10011</b>  | <b>00</b>   | <b>logic.i64 rc,rb,fun3</b> |   |   |   |   |

the *logic* instruction performs a logic operation on the value in the *rb* register then stores the result in the *rc* register. the type of logic operations in *fun3* can be one of:

| value | mnemonic     | description                 |
|-------|--------------|-----------------------------|
| 0     | <i>mv</i>    | <i>move</i>                 |
| 1     | <i>not</i>   | <i>logical not</i>          |
| 2     | <i>neg</i>   | <i>negate</i>               |
| 3     | <i>bswap</i> | <i>bswap</i>                |
| 4     | <i>ctz</i>   | <i>count trailing zeros</i> |
| 5     | <i>clz</i>   | <i>count leading zeros</i>  |
| 6     | <i>ctpop</i> | <i>count population</i>     |
| 7     | <i>sext</i>  | <i>sign extend</i>          |

Table 2.3.: Logic instruction functions

```
match fun3
| mov  -> reg[rc] = reg[rb]
| not  -> reg[rc] = ~reg[rb]
| neg  -> reg[rc] = -reg[rb]
| bswap -> reg[rc] = bswap(reg[rb])
| ctz  -> reg[rc] = ctz(reg[rb])
| clz  -> reg[rc] = clz(reg[rb])
| ctpop -> reg[rc] = ctpop(reg[rb])
| sext -> reg[rc] = sext(reg[rb])
```

### 2.1.21. pin

|           |           |           |               |             |                         |   |   |   |   |
|-----------|-----------|-----------|---------------|-------------|-------------------------|---|---|---|---|
| 15        | 13        | 12        | 10            | 9           | 7                       | 6 | 2 | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> | <i>size</i> |                         |   |   |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>10100</b>  | <b>00</b>   | <b>pin.i64 rc,rb,ra</b> |   |   |   |   |

the *pin* or *pack-indirect* instruction packs two absolute addresses as an *i32x2* (*pc,ib*) relative address vector. the register *ra* is subtracted from the *program counter*, and the register *rb* is subtracted from the *immediate base* register, and the results are packed into an *i32x2* relative address vector and saved to the register *rc*.

```
lpc  = int<i32>(pc - reg[ra])
lib  = int<i32>(ib - reg[rb])
lval = (lpc,lib)
reg[rc] = auth-encrypt(lval)
```



## 2. Instructions

### 2.1.22. and

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>10101</b>  |   |   |   | <b>00</b>   |   |   |

**and.i64** *rc,rb,ra*

the *and* instruction performs a *logical-and* of the register *rb* and the register *ra* and saves the result in the register *rc*.

$$\text{reg}[rc] = \text{reg}[rb] \ \& \ \text{reg}[ra]$$

### 2.1.23. or

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>10110</b>  |   |   |   | <b>00</b>   |   |   |

**or.i64** *rc,rb,ra*

the *or* instruction performs a *logical-or* of the register *rb* and the register *ra* and saves the result in the register *rc*.

$$\text{reg}[rc] = \text{reg}[rb] \ | \ \text{reg}[ra]$$

### 2.1.24. xor

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>10111</b>  |   |   |   | <b>00</b>   |   |   |

**xor.i64** *rc,rb,ra*

the *xor* instruction performs a *logical-exclusive-or* of the register *rb* and the register *ra* and saves the result in the register *rc*.

$$\text{reg}[rc] = \text{reg}[rb] \ \wedge \ \text{reg}[ra]$$

### 2.1.25. add

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>11000</b>  |   |   |   | <b>00</b>   |   |   |

**add.i64** *rc,rb,ra*

the *add* instruction adds the *rb* register to the *ra* register and saves the result in the *rc* register.

$$\text{reg}[rc] = \text{reg}[rb] \ + \ \text{reg}[ra]$$

## 2. Instructions

### 2.1.26. srl

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>11001</b>  |   |   |   | <b>00</b>   |   |   |

**srl.i64** *rc,rb,ra*

the *srl* or *shift-right-logical* instruction performs a logical right shift of the value in the *rb* register by the number of bits in register *ra* then saves the result in the *rc* register. zeros are copied into the right most bits.

$$\text{reg}[rc] = \text{reg}\langle u64 \rangle[rb] \gg \text{reg}[ra]$$

### 2.1.27. sra

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>11010</b>  |   |   |   | <b>00</b>   |   |   |

**sra.i64** *rc,rb,ra*

the *sra* or *shift-right-arithmetic* instruction performs an arithmetic right shift of the value in the *rb* register by the number of bits in register *ra* then saves the result in the *rc* register. sign is copied into the right most bits.

$$\text{reg}[rc] = \text{reg}\langle i64 \rangle[rb] \gg \text{reg}[ra]$$

### 2.1.28. sll

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>11011</b>  |   |   |   | <b>00</b>   |   |   |

**sll.i64** *rc,rb,ra*

the *sll* or *shift-left-logical* instruction performs a logical left shift of the value in the *rb* register by the number of bits in register *ra* then saves the result in the *rc* register.

$$\text{reg}[rc] = \text{reg}[rb] \ll \text{reg}[ra]$$

### 2.1.29. sub

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>11100</b>  |   |   |   | <b>00</b>   |   |   |

**sub.i64** *rc,rb,ra*

the *sub* instruction subtracts the *ra* register from the *rb* register and saves the result in the *rc* register.

$$\text{reg}[rc] = \text{reg}[rb] - \text{reg}[ra]$$

## 2. Instructions

### 2.1.30. mul

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>11101</b>  |   |   |   | <b>00</b>   |   |   |

**mul.i64** *rc,rb,ra*

the *mul* instruction performs signed multiplication of the *rb* register with the *ra* register and saves the result in the *rc* register.

```
reg[rc] = reg[rb] * reg[ra]
```

### 2.1.31. div

|           |           |           |               |   |   |   |             |   |   |
|-----------|-----------|-----------|---------------|---|---|---|-------------|---|---|
| 15        | 13        | 12        | 10            | 9 | 7 | 6 | 2           | 1 | 0 |
| <i>rc</i> | <i>rb</i> | <i>ra</i> | <i>opcode</i> |   |   |   | <i>size</i> |   |   |
| <b>rc</b> | <b>rb</b> | <b>ra</b> | <b>11110</b>  |   |   |   | <b>00</b>   |   |   |

**div.i64** *rc,rb,ra*

the *div* instruction performs signed division of the *rb* register by the *ra* register and saves the result in the *rc* register. division by zero causes the flag to be set and zero to be stored in the *rc* register. a subsequent branch can handle the division by zero.

```
flag = reg[ra] == 0
if flag:
    reg[rc] = 0
else:
    reg[rc] = reg[rb] / reg[ra]
```

### 2.1.32. illegal

|             |   |   |               |             |   |
|-------------|---|---|---------------|-------------|---|
| 15          | 7 | 6 | 2             | 1           | 0 |
| <i>imm9</i> |   |   | <i>opcode</i> | <i>size</i> |   |
| <b>uimm</b> |   |   | <b>11111</b>  | <b>00</b>   |   |

**illegal** *uimm9*

the *illegal* instruction causes an illegal instruction trap. *program counter* and trap cause are saved to privileged registers for the operating system to dispatch to an illegal instruction handler and the *program counter* is set to a trap vector address.

## **3. Assembly**

### **3.1. Introduction**

this glyph assembly language reference begins with an introduction to assembler and linker concepts and terminology, followed by sections describing the glyph assembler directives, and pseudo-instruction aliases. section 2 contains a complete listing of instruction.

### **3.2. Concepts**

this section covers assembler high level concepts required to understand the concepts involved in assembling and linking executable code from source files. the terminology used in this section is applicable to the PE/COFF, ELF and Mach-O file formats.

#### **3.2.1. assembly**

assembly files contains assembly language directives, macros and instructions describing program code and data. they can be handwritten or emitted by a compiler. an assembly file is the input file to the assembler and the output from the assembler is an object file.

#### **3.2.2. object**

object files contain compiled relocatable object code and data emitted by the assembler. an object file cannot be run, rather it is used as input to the linker in a program linking step which combines them to produce an executable file or shared library.

#### **3.2.3. archive**

archive files contain collections of relocatable object files and are typically referred to as static libraries. archive files use a simple format that appends together a set of object files with an index listing the object files contained in the archive. thin archive files just contain the index listing the object files without any associated data.

#### **3.2.4. executable**

executable files contain compiled relocatable object code and data that has been linked together by a linker in a program linking step using multiple object files, archive files and shared libraries as input. executable files can be statically or dynamically linked. dynamically linked executables files express dependencies on shared libraries.

#### 3.2.5. shared library

shared library files contain compiled relocatable object code and data that have been linked together by a linker in a program linking step using multiple object files, archive files and shared libraries as input. shared libraries are dynamically linked by a runtime linker, and they express dependencies on other shared libraries which need to be loaded at runtime.

#### 3.2.6. section

a section is a name for a region of code or data in an object file, executable file or shared library file. the files can be made up of multiple sections where each section corresponds to several types of executable code or data. this list contains the most common types:

- `.text` is a read-only section containing executable code
- `.const` is a read-only section containing immediate blocks
- `.data` is a read-write section containing global or static variables
- `.rodata` is a read-only section containing read-only variables
- `.bss` is a read-write section containing uninitialized data

#### 3.2.7. segment

segments are loadable regions of code or data in an an executable or shared library. segments describe virtual addresses, file offsets and memory access permissions for mapped sections. in ELF, a segment can map to one or more sections. in PE/COFF sections are mapped directly. in Mach-O, sections are contained within a small set of specific segment types.

#### 3.2.8. symbol

symbols are metadata table entries that contains a name with a mapping to an address. symbols are present in object files and dynamic symbols are present in shared libraries and executables. they can be referred to in relocation entries and debugging metadata.

#### 3.2.9. relocation

relocations are metadata table entries used to update relocatable addresses during linking. relocations contains a type, a file offset pointing to text or data, and a pointer to a symbol name whose address needs to be updated during the linking step. relocations can be present in object files and dynamic relocations are present in shared libraries and executables.

#### 3.2.10. program linking

program linking is the process of combining multiple relocatable object files by merging and aligning sections, resolving symbol references across files, assigning final symbol addresses, applying relocation fixups using relocation entries, and adding debug metadata.

### 3. Assembly

## 3.3. Directives

the assembler implements a number of directives that control the assembly of instructions into object files. these directives are based on the AT&T System V assembler [2] and the GNU assembler [5], with some additions. they provide the ability to include arbitrary data, align data, export symbols, switch sections, define constants, and emit metadata.

the following table lists glyph assembler directives:

| Directive                       | Arguments                                            | Description                                                |
|---------------------------------|------------------------------------------------------|------------------------------------------------------------|
| <b>Data directives</b>          |                                                      |                                                            |
| <code>.byte</code>              | <i>expression-list</i>                               | 8-bit comma separated words                                |
| <code>.short</code>             | <i>expression-list</i>                               | 16-bit comma separated words                               |
| <code>.long</code>              | <i>expression-list</i>                               | 32-bit comma separated words                               |
| <code>.quad</code>              | <i>expression-list</i>                               | 64-bit comma separated words                               |
| <code>.octa</code>              | <i>expression-list</i>                               | 128-bit comma separated words                              |
| <code>.string</code>            | <i>"string"</i>                                      | emit string                                                |
| <code>.zero</code>              | <i>integer</i>                                       | emit zeroes                                                |
| <b>Alignment directives</b>     |                                                      |                                                            |
| <code>.align</code>             | <i>pow2 [,pad_val=0] [,max]</i>                      | align to power of 2                                        |
| <code>.balign</code>            | <i>bytes [,pad_val=0]</i>                            | byte align                                                 |
| <b>Symbol directives</b>        |                                                      |                                                            |
| <code>.globl</code>             | <i>symbol_name,const_name</i>                        | emit symbol ( <i>global scope</i> )                        |
| <code>.local</code>             | <i>symbol_name,const_name</i>                        | emit symbol ( <i>local scope</i> )                         |
| <b>Section directives</b>       |                                                      |                                                            |
| <code>.text</code>              | <i>symbol_name,size,align</i><br><i>section_name</i> | emit <code>.text</code> section or make current            |
| <code>.const</code>             |                                                      | emit <code>.const</code> section or make current           |
| <code>.data</code>              |                                                      | emit <code>.data</code> section or make current            |
| <code>.rodata</code>            |                                                      | emit <code>.rodata</code> section or make current          |
| <code>.bss</code>               |                                                      | emit <code>.bss</code> section or make current             |
| <code>.common</code>            |                                                      | emit common object to <code>.bss</code> section            |
| <code>.section</code>           |                                                      | emit section (default <code>.text</code> ) or make current |
| <b>Miscellaneous directives</b> |                                                      |                                                            |
| <code>.equ</code>               | <i>name, value</i>                                   | constant definition                                        |
| <code>.file</code>              | <i>"filename"</i>                                    | emit filename symbol                                       |
| <code>.ident</code>             | <i>"string"</i>                                      | emit identification string                                 |
| <code>.size</code>              | <i>symbol,symbol</i>                                 | emit symbol size                                           |
| <code>.type</code>              | <i>symbol,@function</i>                              | emit symbol type                                           |

Table 3.1.: Assembler directives

### 3. Assembly

## 3.4. Pseudo-instructions

the assembler implements a number of convenience psuedo-instruction aliases that are formed from regular instructions, but have implicit or deduced arguments.

the following table lists glyph assembler pseudo instruction aliases:

| Pseudo-instruction                      | Expansion                                             | Description                         |
|-----------------------------------------|-------------------------------------------------------|-------------------------------------|
| <code>nop</code>                        | <code>or.i64 r0,r0,r0</code>                          | no-operation                        |
| <code>li rc, expression</code>          | (several expansions)                                  | load immediate                      |
| <code>la rc, symbol</code>              | (several expansions)                                  | load address                        |
| <code>call symbol</code>                | <code>jalib.i64 ibcall(text-label,const-label)</code> | procedure call                      |
| <code>ret</code>                        | <code>jtlib.i64 ibret(text-label,const-label)</code>  | procedure return                    |
| <code>jib.i64 ib64(uimm6)</code>        | <code>link.i64 ib64(uimm6), jib</code>                | jump-immmediate-block               |
| <code>jalib.i64 rc, ib64(uimm6)</code>  | <code>link.i64 rc, ib64(uimm6), jalib</code>          | jump-and-link-immmediate-block      |
| <code>jtlib.i64 rc, ib64(uimm6)</code>  | <code>link.i64 rc, ib64(uimm6), jtlib</code>          | jump-to-link-immmediate-block       |
| <code>jalaib.i64 rc, ib64(uimm6)</code> | <code>link.i64 rc, ib64(uimm6), jalaib</code>         | jump-and-link-add-immmediate-block  |
| <code>cmp.lt.i64 rc, rb</code>          | <code>compare.i64 rc, rb, lt</code>                   | compare less than (signed)          |
| <code>cmp.gt.i64 rc, rb</code>          | <code>compare.i64 rb, rc, lt</code>                   | compare greater than (signed)       |
| <code>cmp.le.i64 rc, rb</code>          | <code>compare.i64 rb, rc, ge</code>                   | compare less or equal (signed)      |
| <code>cmp.ge.i64 rc, rb</code>          | <code>compare.i64 rc, rb, ge</code>                   | compare greater or equal (signed)   |
| <code>cmp.eq.i64 rc, rb</code>          | <code>compare.i64 rc, rb, eq</code>                   | compare equal                       |
| <code>cmp.ne.i64 rc, rb</code>          | <code>compare.i64 rc, rb, ne</code>                   | compare not equal                   |
| <code>cmp.ltu.i64 rc, rb</code>         | <code>compare.i64 rc, rb, ltu</code>                  | compare less than (unsigned)        |
| <code>cmp.gtu.i64 rc, rb</code>         | <code>compare.i64 rb, rc, ltu</code>                  | compare greater than (unsigned)     |
| <code>cmp.leu.i64 rc, rb</code>         | <code>compare.i64 rb, rc, geu</code>                  | compare less or equal (unsigned)    |
| <code>cmp.geu.i64 rc, rb</code>         | <code>compare.i64 rc, rb, geu</code>                  | compare greater or equal (unsigned) |
| <code>cmov.i64 rc, rb</code>            | <code>compare.i64 rc, rb, cmov</code>                 | conditional move                    |
| <code>ncmov.i64 rc, rb</code>           | <code>compare.i64 rc, rb, ncmov</code>                | negated conditional move            |
| <code>mov.i64 rc, rb</code>             | <code>logic.i64 rc, rb, mov</code>                    | copy register                       |
| <code>not.i64 rc, rb</code>             | <code>logic.i64 rc, rb, not</code>                    | logical not                         |
| <code>neg.i64 rc, rb</code>             | <code>logic.i64 rc, rb, neg</code>                    | signed negate                       |
| <code>bswap.i64 rc, rb</code>           | <code>logic.i64 rc, rb, bswap</code>                  | byte swap                           |
| <code>ctz.i64 rc, rb</code>             | <code>logic.i64 rc, rb, ctz</code>                    | count trailing zeros                |
| <code>clz.i64 rc, rb</code>             | <code>logic.i64 rc, rb, clz</code>                    | count leading zeroes                |
| <code>ctpop.i64 rc, rb</code>           | <code>logic.i64 rc, rb, ctpop</code>                  | count population                    |
| <code>sext.i64 rc, rb</code>            | <code>logic.i64 rc, rb, sext</code>                   | sign extend                         |

Table 3.2.: Pseudo instructions

## 3.5. Calling convention

### 3.5.1. calling convention — 16-bit

the 16-bit instruction packet, while intended to be used in conjunction with the larger opcodes, is designed as a complete subset, so there is an ABI variant that targets a subset of the instruction set architecture that only uses the 16-bit opcodes.

the register assignment for the 16-bit subset was chosen with this rationale:

- 2 blocks of 4 contiguous non-volatile *callee-save* and volatile *caller-save* registers.
- 3 special registers, 2 argument registers, 1 temporary register, and 3 save registers.
- 3 save registers to avoid excessive spilling around function calls.
- 1 temporary register to avoid spilling arguments to free a temporary.

the calling convention for the 16-bit subset is as follows:

- *immediate base* `ib` is set by `call` instructions and must point to a valid immediate block on function entry. function symbols are exported with two labels; one in the `.text` section, and one in the `.const` section. *immediate base* must be restored to the entry value in the function *epilogue* before it can be restored by `ret`.
- *argument registers* `a0` and `a1` are used for the first two arguments, and the remaining arguments are passed on the stack. *return value* is places in `a0` and `a1`, *temporary register* `t0` is a volatile register, and *frame pointer* (if enabled) uses `s0`. there are two more non-volatile *callee-save* registers, `s1` and `s2`.

the following table outlines the 16-bit register allocation showing register name alias, description, and non-volatile *callee-save* or volatile *caller-save* status.

| name            | alias              | description                          | save   |
|-----------------|--------------------|--------------------------------------|--------|
| <code>ib</code> |                    | immediate base                       | callee |
| <code>r0</code> | <code>sp</code>    | stack pointer                        | callee |
| <code>r1</code> | <code>s0/fp</code> | saved register 0 / frame pointer     | callee |
| <code>r2</code> | <code>s1</code>    | saved register 1                     | callee |
| <code>r3</code> | <code>s2</code>    | saved register 2                     | callee |
| <code>r4</code> | <code>a0</code>    | argument register 0                  | caller |
| <code>r5</code> | <code>a1</code>    | argument register 1                  | caller |
| <code>r6</code> | <code>t0</code>    | temporary register 0                 | caller |
| <code>r7</code> | <code>ra</code>    | return address / (pc,ib) link vector | caller |

Table 3.3.: 16-bit register assignment



# A. Appendix

## A.1. Opcode summary — 16-bit

| 15      | 13   | 12   | 10 | 9 | 7 | 6      | 2 | 1    | 0 |                                                                          |
|---------|------|------|----|---|---|--------|---|------|---|--------------------------------------------------------------------------|
| operand |      |      |    |   |   | opcode |   | size |   |                                                                          |
| uimm    |      |      |    |   |   | 00000  |   | 00   |   | <b>break</b> <i>uimm9</i>                                                |
| simm    |      |      |    |   |   | 00001  |   | 00   |   | <b>j</b> <i>simm9</i> × 2                                                |
| simm    |      |      |    |   |   | 00010  |   | 00   |   | <b>b</b> <i>simm9</i> × 2                                                |
| simm    |      |      |    |   |   | 00011  |   | 00   |   | <b>ibj</b> <i>simm9</i> × 64                                             |
| fun     | uimm |      |    |   |   | 00100  |   | 00   |   | <b>link.i64</b> <i>fun3</i> , <i>ib64</i> ( <i>uimm6</i> )               |
| rc      | uimm |      |    |   |   | 00101  |   | 00   |   | <b>movh.i64</b> <i>rc</i> , <i>ib32</i> ( <i>uimm6</i> )                 |
| rc      | uimm |      |    |   |   | 00110  |   | 00   |   | <b>movw.i64</b> <i>rc</i> , <i>ib64</i> ( <i>uimm6</i> )                 |
| rc      | simm |      |    |   |   | 00111  |   | 00   |   | <b>movi.i64</b> <i>rc</i> , <i>simm6</i>                                 |
| rc      | simm |      |    |   |   | 01000  |   | 00   |   | <b>addi.i64</b> <i>rc</i> , <i>simm6</i>                                 |
| rc      | uimm |      |    |   |   | 01001  |   | 00   |   | <b>srli.i64</b> <i>rc</i> , <i>uimm6</i>                                 |
| rc      | uimm |      |    |   |   | 01010  |   | 00   |   | <b>srai.i64</b> <i>rc</i> , <i>uimm6</i>                                 |
| rc      | uimm |      |    |   |   | 01011  |   | 00   |   | <b>slli.i64</b> <i>rc</i> , <i>uimm6</i>                                 |
| rc      | uimm |      |    |   |   | 01100  |   | 00   |   | <b>addh.i64</b> <i>rc</i> , <i>ib32</i> ( <i>uimm6</i> )                 |
| rc      | uimm |      |    |   |   | 01101  |   | 00   |   | <b>leapc.i64</b> <i>rc</i> , <i>ib32</i> ( <i>uimm6</i> )( <i>pc</i> )   |
| rc      | uimm |      |    |   |   | 01110  |   | 00   |   | <b>loadpc.i64</b> <i>rc</i> , <i>ib32</i> ( <i>uimm6</i> )( <i>pc</i> )  |
| rc      | uimm |      |    |   |   | 01111  |   | 00   |   | <b>storepc.i64</b> <i>rc</i> , <i>ib32</i> ( <i>uimm6</i> )( <i>pc</i> ) |
| rc      | rb   | uimm |    |   |   | 10000  |   | 00   |   | <b>load.i64</b> <i>rc</i> , ( <i>uimm3</i> × 8)( <i>rb</i> )             |
| rc      | rb   | uimm |    |   |   | 10001  |   | 00   |   | <b>store.i64</b> <i>rc</i> , ( <i>uimm3</i> × 8)( <i>rb</i> )            |
| rc      | rb   | fun  |    |   |   | 10010  |   | 00   |   | <b>compare.i64</b> <i>rc</i> , <i>rb</i> , <i>fun3</i>                   |
| rc      | rb   | fun  |    |   |   | 10011  |   | 00   |   | <b>logic.i64</b> <i>rc</i> , <i>rb</i> , <i>fun3</i>                     |
| rc      | rb   | ra   |    |   |   | 10100  |   | 00   |   | <b>pin.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| rc      | rb   | ra   |    |   |   | 10101  |   | 00   |   | <b>and.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| rc      | rb   | ra   |    |   |   | 10110  |   | 00   |   | <b>or.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                          |
| rc      | rb   | ra   |    |   |   | 10111  |   | 00   |   | <b>xor.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| rc      | rb   | ra   |    |   |   | 11000  |   | 00   |   | <b>add.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| rc      | rb   | ra   |    |   |   | 11001  |   | 00   |   | <b>srl.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| rc      | rb   | ra   |    |   |   | 11010  |   | 00   |   | <b>sra.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| rc      | rb   | ra   |    |   |   | 11011  |   | 00   |   | <b>sll.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| rc      | rb   | ra   |    |   |   | 11100  |   | 00   |   | <b>sub.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| rc      | rb   | ra   |    |   |   | 11101  |   | 00   |   | <b>mul.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| rc      | rb   | ra   |    |   |   | 11110  |   | 00   |   | <b>div.i64</b> <i>rc</i> , <i>rb</i> , <i>ra</i>                         |
| uimm    |      |      |    |   |   | 11111  |   | 00   |   | <b>illegal</b> <i>uimm9</i>                                              |

# References

- [1] Howard H. Aiken and Grace Hopper. *A Manual of Operation for the Automatic Sequence Controlled Calculator*. Tech. rep. Describes the Harvard Mark I architecture with dedicated instruction memory. Harvard University Computation Laboratory, 1946.
- [2] AT&T Bell Laboratories. *UNIX<sup>®</sup> System V Release 4: Assembler and Machine-Level Debugging Guide*. Describes AT&T assembler directives and syntax used in UNIX<sup>®</sup> System V Release 4. Prentice Hall, 1990. ISBN: 0-13-947116-4.
- [3] J. Cocke and V. Markstein. “The Evolution of RISC Technology at IBM”. In: *IBM Journal of Research and Development* 34.1 (1990). Reviews the IBM 801’s architecture and its influence on later RISC systems, pp. 4–11. DOI: 10.1147/rd.341.0004.
- [4] J. Cocke and V. Markstein. “The Evolution of the RISC System/370”. In: *IBM Conference on Directions in Computer Architecture* (1980). Describes the IBM 801 architecture, an early example of a load-store RISC design, pp. 49–50.
- [5] Free Software Foundation. *Using as: The GNU Assembler*. Part of GNU Binutils. GNU Project. 1991. URL: <https://sourceware.org/binutils/docs/as/>.
- [6] Advanced Micro Devices Inc. “Accessing an extended register set in an extended register mode”. Patent US6877084B1. Filed 2001-04-02. Granted 2005-04-05. Expired 2023-06-18. Apr. 2005. URL: <https://patents.google.com/patent/US6877084B1>.
- [7] Cray Research LLC. “Data processing system for processing one and two parcel instructions”. Patent US5717881A. Filed 1995-06-07. Granted 1998-02-10. Expired 2015-02-10. Feb. 1998. URL: <https://patents.google.com/patent/US5717881A>.
- [8] John W. Mauchly and J. Presper Eckert. *Description of the ENIAC and Comments on Electronic Digital Computing Machines*. Tech. rep. Describes the ENIAC architecture with a unified instruction and data path, an early example of a Von Neumann-style design. Moore School of Electrical Engineering, University of Pennsylvania, 1946.
- [9] John von Neumann. *First Draft of a Report on the EDVAC*. Tech. rep. Seminal report introducing the stored-program concept, later known as the Von Neumann architecture. Moore School of Electrical Engineering, University of Pennsylvania, 1945.
- [10] David A. Patterson and John L. Hennessy. *Computer Organization and Design: The Hardware Software Interface*. 5th. Describes the classic 5-stage RISC pipeline: fetch, decode, execute, memory, write-back. Morgan Kaufmann, 2013. ISBN: 9780124077263.
- [11] Ray J. Solomonoff. “A Formal Theory of Inductive Inference. Parts I and II”. In: *Information and Control* 7.1-2 (1964). Introduces the three-tape Universal Prefix Turing Machine model, pp. 1–22, 224–254. DOI: 10.1016/S0019-9958(64)90223-2.
- [12] Alan M. Turing. “On Computable Numbers, with an Application to the Entscheidungsproblem”. In: *Proceedings of the London Mathematical Society*. 2nd ser. 42 (1936). A theoretical model of computation, pp. 230–265. DOI: 10.1112/plms/s2-42.1.230.