

GLYPH-X

a super regular RISC architecture

Michael Clark

April 6, 2025 DRAFT

Contents

1. Architecture	2
1.1. Introduction	2
1.2. Instruction format	3
1.2.1. instruction templates	3
1.2.2. size encoding	4
1.2.3. 16-bit instructions	4
1.3. Constant stream	5
1.4. Register file	6
2. Instructions	7
2.1. 16-bit opcodes	7
2.1.1. break	7
2.1.2. j	7
2.1.3. b	7
2.1.4. ibj	8
2.1.5. jalib.i64	8
2.1.6. jtlib.i64	8
2.1.7. movib.i64	9
2.1.8. movi.i64	9
2.1.9. addi.i64	9
2.1.10. srli.i64	9
2.1.11. srai.i64	10
2.1.12. slli.i64	10
2.1.13. addib.i64	10
2.1.14. leapc.i64	10
2.1.15. loadpc.i64	11
2.1.16. storepc.i64	11
2.1.17. load.i64	11
2.1.18. store.i64	11
2.1.19. compare.i64	12
2.1.20. logic.i64	12
2.1.21. pin.i64	13
2.1.22. and.i64	13
2.1.23. or.i64	13
2.1.24. xor.i64	13
2.1.25. add.i64	14
2.1.26. srl.i64	14

Contents

2.1.27. sra.i64	14
2.1.28. sll.i64	14
2.1.29. sub.i64	15
2.1.30. mul.i64	15
2.1.31. div.i64	15
2.1.32. illegal	15
3. Assembler	16
3.1. Introduction	16
3.2. Concepts	16
3.2.1. assembly file	16
3.2.2. relocatable object file	16
3.2.3. file header	16
3.2.4. program header	17
3.2.5. section header	17
3.2.6. sections	17
3.2.7. program linking	17
3.2.8. linker script	17
3.3. Directives	18
3.4. Pseudo-instructions	19
3.5. Calling convention	20
3.5.1. 16-bit calling convention	20
A. Appendix	21
A.1. 16-bit opcodes	22

1. Architecture

1.1. Introduction

glyph is a super regular RISC architecture that encodes constants in a secondary stream accessed via an immediate base register that points at immediate blocks containing constants accessed via a constant address mode. the immediate base register branches like the program counter, and procedure call and return set and restore (pc, ib) together.

glyph uses relative address vectors in its link register which is different to typical RISC architectures. glyph does this so that the branch instructions can fit (pc, ib) into the a single link register for compatibility. glyph achieves this by packing two relative (pc, ib) displacements into a relative address vector.

immediate blocks can be switched using the immediate block branch instruction. immediate blocks, unlike typical RISC architectures, mean that most relocations are word sized like CISC architectures, and can use C-style structure packing and alignment rules.

this list outlines some differentiating elements of the super regular RISC architecture:

- variable length instruction format supporting 16, 32, 64, and 128-bit instructions.
- 16-bit compressed instruction packets can access 8 registers.
- (pc, ib) is a special program counter and immediate base register address vector.
- link register contains packed relative (pc, ib) address vector to function entry.
- `ibj` *immediate-block-jump* adds a relative address to the immediate base register.
- `lib` *load-immediate-block* uses an unsigned displacement to access constants.
- `jlib` *jump-and-link-immediate-block* or *call* links address vector and adds constants to (pc, ib) and is used to branch the program counter and immediate base register at the same time for calling procedures.
- `jtlib` *jump-to-link-immediate-block* or *ret* subtracts link vector from and adds constants to (pc, ib) and is used to branch the program counter and immediate base register at the same time for returning from called procedures.
- `pin` *pack-indirect* packs two absolute addresses as relative address vector from (pc, ib) and is used for calling absolute addresses such as virtual functions.

1. Architecture

1.2. Instruction format

glyph has a variable length instruction format supporting 16, 32, 64, and 128-bit instruction packets. the instruction packet has been designed to use a *super regular* scheme, whereby successive instruction packets extend the fields in the previous packet.

1.2.1. instruction templates

the variable length instruction format has a single base format where fields in the template instruction form are extended by successive instruction packets.

15	7	6	2	1	0
<i>operand</i> _[8:0]		<i>opcode</i> _[4:0]		<i>sz</i> _[1:0]	

Figure 1.1.: 16-bit instruction template.

15	7	6	2	1	0
<i>operand</i> _[8:0]		<i>opcode</i> _[4:0]		<i>sz</i> _[1:0]	
<i>operand</i> _[17:9]		<i>opcode</i> _[9:5]		<i>sz</i> _[3:2]	

Figure 1.2.: 32-bit instruction template.

15	7	6	2	1	0
<i>operand</i> _[8:0]		<i>opcode</i> _[4:0]		<i>sz</i> _[1:0]	
<i>operand</i> _[17:9]		<i>opcode</i> _[9:5]		<i>sz</i> _[3:2]	
<i>operand</i> _[26:18]		<i>opcode</i> _[14:10]		<i>sz</i> _[5:4]	
<i>operand</i> _[35:27]		<i>opcode</i> _[19:15]		<i>sz</i> _[7:6]	

Figure 1.3.: 64-bit instruction template.

15	7	6	2	1	0
<i>operand</i> _[8:0]		<i>opcode</i> _[4:0]		<i>sz</i> _[1:0]	
<i>operand</i> _[17:9]		<i>opcode</i> _[9:5]		<i>sz</i> _[3:2]	
<i>operand</i> _[26:18]		<i>opcode</i> _[14:10]		<i>sz</i> _[5:4]	
<i>operand</i> _[35:27]		<i>opcode</i> _[19:15]		<i>sz</i> _[7:6]	
<i>operand</i> _[44:36]		<i>opcode</i> _[24:20]		<i>sz</i> _[9:8]	
<i>operand</i> _[53:45]		<i>opcode</i> _[29:25]		<i>sz</i> _[11:10]	
<i>operand</i> _[62:54]		<i>opcode</i> _[34:30]		<i>sz</i> _[13:12]	
<i>operand</i> _[71:63]		<i>opcode</i> _[39:35]		<i>sz</i> _[15:14]	

Figure 1.4.: 128-bit instruction template.

1. Architecture

1.2.2. size encoding

the variable length instruction format has a *2-bit* size field in a fixed position in every 16-bit instruction packet, somewhat inspired by LEB128, to reduce the complexity of variable length instruction size decoding.

Instruction Size	Size Fields
16-bit	{00}
32-bit	{01,11}
64-bit	{10,11,11,11}
128-bit	{11,11,11,11,11,11,11,11}

Table 1.1.: Variable-length instruction size fields

1.2.3. 16-bit instructions

the 16-bit instructions forms are super regular in that operand and opcode bits do not overlap and the number and complexity of the formats is reduced so that vectorized instruction decoding is easier in software. the scheme is designed so that *1-bit* of coding space in the larger packet can be used to extend register sizes for the 16-bit ops.

15	7	6	2	1	0
<i>imm</i> _[5:0]			<i>opcode</i> _[4:0]	00	

Figure 1.5.: 16-bit large immediate.

15	13	12	7	6	2	1	0
<i>rc</i> _[2:0]		<i>imm</i> _[5:0]		<i>opcode</i> _[4:0]	00		

Figure 1.6.: 16-bit one operand with immediate.

15	13	12	10	9	7	6	2	1	0
<i>rc</i> _[2:0]		<i>rb</i> _[2:0]		<i>imm</i> _[2:0]		<i>opcode</i> _[4:0]	00		

Figure 1.7.: 16-bit two operand with immediate.

15	13	12	10	9	7	6	2	1	0
<i>rc</i> _[2:0]		<i>rb</i> _[2:0]		<i>ra</i> _[2:0]		<i>opcode</i> _[4:0]	00		

Figure 1.8.: 16-bit three operand.

1.3. Constant stream

glyph separates the instruction stream into two streams, one with instructions and one with constants. the instruction stream is addressed with the *program counter* (pc) and the constant stream is addressed with the *immediate base register* (ib).




Figure 1.9.: program counter and immediate base register.

immediate blocks are aligned memory blocks addressed by the *immediate base register*.

immediate blocks can be navigated by branching the constant stream independently using the *immediate-block-jump* instruction, or together with the program counter using procedure call and return instructions that set and restore (pc, ib) via a link register that contains a packed relative address vector. the use of packed relative address vectors is for backward compatibility with a single link register.

for procedure calls and returns, the instruction and constant streams are set at the same time using the call and return instructions; *jump-and-link-immediate-block*, *jump-to-link-immediate-block*, which add and subtract relative address vectors to (pc, ib) , the program counter and the immediate base register. the *pack-indirect* instruction allows absolute (pc, ib) addresses to be packed into a relative address vector for indirect calls.

the instruction forms only use bonded register slots for immediate operands and operand bits do not overlap opcode bits. the use of immediate blocks means large immediate constants can all be accessed with short references encoded inside of register slots for instructions that use an immediate block relative addressing mode.

1.4. Register file

the glyph register file is extensible due to the variable length instruction format and supports a different number of registers depending on the instruction size.

- 16-bit instruction packet can access 8 registers with up to 3 operands.
- 32-bit instruction packet can access 64 registers with up to 3 operands.
- 64-bit instruction packet can access 64 registers with up to 6 operands.

the register state accessible by the 16-bit instruction packet is comprised of:

- program counter register (aligned to 2 bytes).
- immediate base register (aligned to 64 bytes).
- 8 × general purpose registers.
- 1 × predicate register.

the following diagram shows the register state accessible by the 16-bit instruction packet:




Figure 1.10.: register state accessible by 16-bit instruction packet.

2. Instructions

2.1. 16-bit opcodes

2.1.1. break

15	7	6	2	1	0
<i>imm9</i>			<i>opcode</i>	<i>size</i>	
uimm			00000	00	

break uimm9

The *break* instruction causes a trap to the debugger. The program counter and trap cause are saved to privileged registers for the operating system to dispatch to a debugger.

2.1.2. j

15	7	6	2	1	0
<i>imm9</i>			<i>opcode</i>	<i>size</i>	
simm			00001	00	

j simm9

The *j* or *jump* instruction is an unconditional branch instruction that adds a relative immediate address to the program counter. The resulting address is $pc + \text{simm9} + 2$.

```
pc += simm9 * 2 + 2
```

2.1.3. b

15	7	6	2	1	0
<i>imm9</i>			<i>opcode</i>	<i>size</i>	
simm			00010	00	

b simm9

The *b* or *branch* instruction is a conditional branch instruction that adds a relative immediate address to the program counter, if the *flag* register has been set by a compare instruction. The resulting address is $pc + \text{simm9} + 2$.

```
if flag:
    pc += simm9 * 2 + 2
```

2. Instructions

2.1.4. ibj

15	7	6	2	1	0
<i>imm9</i>			<i>opcode</i>		<i>size</i>
simm			00011		00

ibj *simm9*

the *ibj* or *immediate-block-jump* instruction adds a 64-bit relative address to the immediate base register.

```
ib += simm9 * 64
```

2.1.5. jalib.i64

15	13	12	7	6	2	1	0
<i>rc</i>		<i>imm6</i>			<i>opcode</i>		<i>size</i>
rc		uimm			00100		00

jalib *rc*, *ib64*(*uimm6**8)

the *jalib* or *jump-and-link-immediate-block* instruction loads a 64-bit constant addressed by [*ib* + *uimm6* * 8] containing a *i32x2* relative address vector, adds it to (*pc*,*ib*), then saves the relative address vector to the *rc* register.

```
tmp = (i32x2)[ib + uimm6 * 8]
{rpc,rib} = (i32x2)tmp
pc += rpc + 2
ib += rib
r[rc] = tmp
```

2.1.6. jtlib.i64

15	13	12	7	6	2	1	0
<i>rc</i>		<i>imm6</i>			<i>opcode</i>		<i>size</i>
rc		uimm			00101		00

jtlib *rc*, *ib64*(*uimm6**8)

the *jtlib* or *jump-to-link-immediate-block* instruction loads a 64-bit constant addressed by [*ib* + *uimm6* * 8] containing a *i32x2* relative address vector, subtracts the *i32x2* relative address vector in the *rc* register from it, then adds it to (*pc*,*ib*).

```
{rpc,rib} = (i32x2)r[rc]
{dpc,dib} = (i32x2)[ib + uimm6 * 8]
pc += dpc - rpc
ib += dib - rib
```

2. Instructions

2.1.7. movib.i64

15	13	12	7	6	2	1	0
<i>rc</i>	<i>imm6</i>			<i>opcode</i>	<i>size</i>		
rc	uimm			00110	00		

movib.i64 rc, ib64(uimm6*8)

the *movib* or *move-immediate-block* instruction loads a 64-bit constant addressed by $[ib + uimm6 * 8]$ then saves it to the *rc* register.

```
r[rc] = (i64)[ib + uimm6 * 8]
```

2.1.8. movi.i64

15	13	12	7	6	2	1	0
<i>rc</i>	<i>imm6</i>			<i>opcode</i>	<i>size</i>		
rc	simm			00111	00		

movi.i64 rc, simm6

the *movi* or *move-immediate* instruction sign-extends the immediate value in *simm6* then saves it to the *rc* register.

```
r[rc] = simm6
```

2.1.9. addi.i64

15	13	12	7	6	2	1	0
<i>rc</i>	<i>imm6</i>			<i>opcode</i>	<i>size</i>		
rc	simm			01000	00		

addi.i64 rc, simm6

the *addi* or *add-immediate* instruction sign-extends the immediate value in *simm6* then adds it to the *rc* register.

```
r[rc] = r[rc] + simm6
```

2.1.10. srli.i64

15	13	12	7	6	2	1	0
<i>rc</i>	<i>imm6</i>			<i>opcode</i>	<i>size</i>		
rc	uimm			01001	00		

srli.i64 rc, uimm6

the *srli* or *shift-right-logical-immediate* instruction performs a logical right shift by *uimm6* bits of the value in the *rc* register. zeros are copied into the right most bits.

```
r[rc] = (u64)r[rc] >> uimm6
```

2. Instructions

2.1.11. srai.i64

15	13	12	7	6	2	1	0
<i>rc</i>	<i>imm6</i>			<i>opcode</i>	<i>size</i>		
rc	uimm			01010	00		

srai.i64 rc, uimm6

the *srai* or *shift-right-arithmetic-immediate* instruction performs an arithmetic right shift by *uimm6* bits of the value in the *rc* register. sign is copied into the right most bits.

```
r[rc] = (i64)r[rc] >> uimm6
```

2.1.12. slli.i64

15	13	12	7	6	2	1	0
<i>rc</i>	<i>imm6</i>			<i>opcode</i>	<i>size</i>		
rc	uimm			01011	00		

slli.i64 rc, uimm6

the *slli* or *shift-left-logical-immediate* instruction performs a logical left shift by *uimm6* bits of the value in the *rc* register.

```
r[rc] = r[rc] << uimm6
```

2.1.13. addib.i64

15	13	12	7	6	2	1	0
<i>rc</i>	<i>imm6</i>			<i>opcode</i>	<i>size</i>		
rc	uimm			01100	00		

addib.i64 rc, ib32(uimm6*4)

the *addib* or *add-immediate-block-constant* instruction loads a 32-bit constant addressed by $[ib + uimm6 * 4]$, sign-extends it to 64-bits, adds it to the *rc* register and saves the result in the *rc* register.

```
r[rc] = r[rc] + (i64)[ib + uimm6 * 4]
```

2.1.14. leapc.i64

15	13	12	7	6	2	1	0
<i>rc</i>	<i>imm6</i>			<i>opcode</i>	<i>size</i>		
rc	uimm			01101	00		

leapc.i64 rc, ib32(uimm6*4)(pc)

the *leapc* or *load-effective-address-pc* instruction loads a 32-bit constant addressed by $[ib + uimm6 * 4]$, sign-extends it to 64-bits, adds it to the *program counter* and saves the result in the *rc* register.

```
r[rc] = pc + (i32)[ib + uimm6 * 4]
```

2. Instructions

2.1.15. loadpc.i64

15	13	12	7	6	2	1	0
<i>rc</i>		<i>imm6</i>			<i>opcode</i>		<i>size</i>
rc		uimm			01110		00

loadpc.i64 rc, ib32(uimm6*4)(pc)

the *loadpc* instruction loads a 32-bit constant addressed by $[ib + uimm6 * 4]$ which it sign-extends to 64-bit, adds it to the *program counter* to form an address, then loads a 64-bit value from memory at that address and saves the result in the *rc* register.

$$r[rc] = (i64)[pc + (i32)[ib + uimm6 * 4]]$$

2.1.16. storepc.i64

15	13	12	7	6	2	1	0
<i>rc</i>		<i>imm6</i>			<i>opcode</i>		<i>size</i>
rc		uimm			01111		00

storepc.i64 rc, ib32(uimm6*4)(pc)

the *storepc* instruction loads a 32-bit constant addressed by $[ib + uimm6 * 4]$ which it sign-extends to 64-bit, adds it to the *program counter* to form an address, then stores to memory at that address a 64-bit value from the *rc* register.

$$(i64)[pc + (i64)[ib + uimm6 * 4]] = r[rc]$$

2.1.17. load.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>imm3</i>			<i>opcode</i>		<i>size</i>
rc		rb		uimm			10000		00

load.i64 rc, (uimm3*8)(rb)

the *load* instruction adds $uimm3 * 8$ to the *rb* to form an address, then loads a 64-bit value from memory at that address and saves the result in the *rc* register.

$$r[rc] = (i64)[r[rb] + uimm3 * 8]$$

2.1.18. store.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>imm3</i>			<i>opcode</i>		<i>size</i>
rc		rb		uimm			10001		00

store.i64 rc, (uimm3*8)(rb)

the *store* instruction adds $uimm3 * 8$ to the *rb* register to form an address, then stores to memory at that address a 64-bit value from the *rc* register.

$$(i64)[r[rb] + uimm3 * 8] = r[rc]$$

2. Instructions

2.1.19. compare.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>imm3</i>		<i>opcode</i>		<i>size</i>	
rc		rb		uimm		10010		00	

cmp.i64 rc, rb, fun3

the *compare* instruction performs a comparison between the value in *rb* and *rc* then saves the result to the *flag* register. the *compare* opcode is also used to perform conditional move whereby the *rb* register is copied into the *rc* if the *flag* register is set. the type of comparison in *fun3* can be one of: 0 → *less than (signed)*, 1 → *greater or equal (signed)*, 2 → *equal*, 3 → *not equal*, 4 → *less than (unsigned)*, 5 → *greater or equal (unsigned)*, or 6 → *conditional move*.

```
match fun3
| lt  -> flag = (i64)r[rc] < (i64)r[rb]
| ge  -> flag = (i64)r[rc] >= (i64)r[rb]
| eq  -> flag =      r[rc] =      r[rb]
| ne  -> flag =      r[rc] !=      r[rb]
| ltu -> flag = (u64)r[rc] < (u64)r[rb]
| geu -> flag = (u64)r[rc] >= (u64)r[rb]
| mov -> if (flag) r[rc] = r[rb]
```

2.1.20. logic.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>imm3</i>		<i>opcode</i>		<i>size</i>	
rc		rb		uimm		10011		00	

logic.i64 rc, rb, fun3

the *logic* instruction performs a logic operation on the value in the *rb* register then stores the result in the *rc* register. the type of logic operations in *fun3* can be one of: 0 → *move*, 1 → *logical not*, 2 → *negate*, 3 → *bswap*, 4 → *count trailing zeros*, 5 → *count leading zeros*, or 6 → *count population*.

```
match fun3
| mov  -> r[rc] = r[rb]
| not  -> r[rc] = ~r[rb]
| neg  -> r[rc] = -r[rb]
| bswap -> r[rc] = bswap(r[rb])
| ctz  -> r[rc] = ctz(r[rb])
| clz  -> r[rc] = clz(r[rb])
| ctpop -> r[rc] = ctpop(r[rb])
```

2. Instructions

2.1.21. pin.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>ra</i>					<i>size</i>
rc		rb		ra			10100		00

pin.i64 rc, rb, ra

the *pin* or *pack-indirect* instruction packs two absolute addresses as an *i32x2* (*pc,ib*) relative address vector. the register *ra* is subtracted from the *program counter + 2*, and the register *rb* is subtracted from *immediate base register*, and the results are packed into an *i32x2* relative address vector and saved to the register *rc*.

```
rpc = (i32)(pc - r[ra] + 2)
rib = (i32)(ib - r[rb]    )
r[rc] = (i32x2){rpc,rib}
```

2.1.22. and.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>ra</i>					<i>size</i>
rc		rb		ra			10101		00

and.i64 rc, rb, ra

the *and* instruction performs a *logical-and* of the register *rb* and the register *ra* and saves the result in the register *rc*.

```
r[rc] = r[rb] & r[ra]
```

2.1.23. or.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>ra</i>					<i>size</i>
rc		rb		ra			10110		00

or.i64 rc, rb, ra

the *or* instruction performs a *logical-or* of the register *rb* and the register *ra* and saves the result in the register *rc*.

```
r[rc] = r[rb] | r[ra]
```

2.1.24. xor.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>ra</i>					<i>size</i>
rc		rb		ra			10111		00

xor.i64 rc, rb, ra

the *xor* instruction performs a *logical-exclusive-or* of the register *rb* and the register *ra* and saves the result in the register *rc*.

```
r[rc] = r[rb] ^ r[ra]
```

2. Instructions

2.1.25. add.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>	<i>rb</i>	<i>ra</i>	<i>opcode</i>				<i>size</i>		
rc	rb	ra	11000				00		

add.i64 rc, rb, ra

the *add* instruction adds the *rb* register to the *ra* register and saves the result in the *rc* register.

$$r[rc] = r[rb] + r[ra]$$

2.1.26. srl.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>	<i>rb</i>	<i>ra</i>	<i>opcode</i>				<i>size</i>		
rc	rb	ra	11001				00		

srl.i64 rc, rb, ra

the *srl* or *shift-right-logical* instruction performs a logical right shift of the value in the *rb* register by the number of bits in register *ra* then saves the result in the *rc* register. zeros are copied into the right most bits.

$$r[rc] = (u64)r[rb] \gg r[ra]$$

2.1.27. sra.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>	<i>rb</i>	<i>ra</i>	<i>opcode</i>				<i>size</i>		
rc	rb	ra	11010				00		

sra.i64 rc, rb, ra

the *sra* or *shift-right-arithmetic* instruction performs an arithmetic right shift of the value in the *rb* register by the number of bits in register *ra* then saves the result in the *rc* register. sign is copied into the right most bits.

$$r[rc] = (i64)r[rb] \gg r[ra]$$

2.1.28. sll.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>	<i>rb</i>	<i>ra</i>	<i>opcode</i>				<i>size</i>		
rc	rb	ra	11011				00		

sll.i64 rc, rb, ra

the *sll* or *shift-left-logical* instruction performs a logical left shift of the value in the *rb* register by the number of bits in register *ra* then saves the result in the *rc* register.

$$r[rc] = r[rb] \ll r[ra]$$

2. Instructions

2.1.29. sub.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>ra</i>			<i>opcode</i>		<i>size</i>
rc		rb		ra			11100		00

sub.i64 rc, rb, ra

the *sub* instruction subtracts the *ra* register from the *rb* register and saves the result in the *rc* register.

$$r[rc] = r[rb] - r[ra]$$

2.1.30. mul.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>ra</i>			<i>opcode</i>		<i>size</i>
rc		rb		ra			11101		00

mul.i64 rc, rb, ra

the *mul* instruction performs signed multiplication of the *rb* register with the *ra* register and saves the result in the *rc* register.

$$r[rc] = r[rb] * r[ra]$$

2.1.31. div.i64

15	13	12	10	9	7	6	2	1	0
<i>rc</i>		<i>rb</i>		<i>ra</i>			<i>opcode</i>		<i>size</i>
rc		rb		ra			11110		00

div.i64 rc, rb, ra

the *div* instruction performs signed division of the *rb* register by the *ra* register and saves the result in the *rc* register. division by zero causes a divide by zero exception.

$$r[rc] = r[rb] / r[ra]$$

2.1.32. illegal

15	7	6	2	1	0
<i>imm9</i>			<i>opcode</i>		<i>size</i>
uimm			11111		00

illegal uimm9

the *illegal* instruction causes an illegal instruction trap.

3. Assembler

3.1. Introduction

this glyph assembly language reference begins with an introduction to assembler and linker concepts, followed by sections describing the glyph assembler directives, and pseudo-instruction aliases. section 2 contains a complete listing of instruction.

3.2. Concepts

this section covers assembler high level concepts required to understand the concepts involved in assembling and linking executable code from source files.

3.2.1. assembly file

an assembly file contains assembly language directives, macros and instructions. it can be emitted by a compiler or it can be handwritten. an assembly file is the input file to the assembler. the extensions for assembly files are `.s`.

3.2.2. relocatable object file

relocatable object files contain compiled object code and data emitted by the assembler. an object file cannot be run, rather it is used as input to the linker as a step towards producing an executable file. the extension for object files is `.o`.

3.2.3. file header

an assembler file has a file header that contains magic to indicate how the file is formatted, the architecture of the binary, the endianness of the binary; *little-endian* in the case of glyph, the file type (*relocatable object*, *executable*, *shared library*), the number of program headers and their offsets in the file, the number of section headers and their offsets in the file, plus fields indicating the file format version and various other details.

3.2.4. program header

program headers provide size and offsets of loadable segments within an executable or shared library along with protection attributes used by the operating system (*read, write and exec*). program headers are not present in relocatable object files and are primarily for use by the operating system to and dynamic linker to map code and data into memory.

3.2.5. section header

section headers provide size, offset, type, alignment and flags for the sections contained within the binary file. section headers are not required to execute a static binary but are necessary for dynamic linking and program linking. various section types refer to the location of the symbol table, relocations and dynamic symbols in the binary file.

3.2.6. sections

an object file is made up of multiple sections, with each section corresponding to distinct types of executable code or data. there are a variety of different section types. this list contains the four most common sections:

- `.text` is a read-only section containing executable code
- `.const` is a read-only section containing immediate blocks
- `.data` is a read-write section containing global or static variables
- `.rodata` is a read-only section containing read-only variables
- `.bss` is a read-write section containing uninitialized data

3.2.7. program linking

program linking is the process of reading multiple relocatable object files, merging the sections from each of the source files, calculating the new addresses for symbols and applying relocation fixups to text or data that is pointed to in relocation entries.

3.2.8. linker script

linker scripts are text source files that are optionally input to the linker containing rules for the linker to use when calculating the load address and alignment of the various sections when creating an executable output file. the extension for linker scripts is `.ld`.

3.3. Directives

the assembler implements a number of directives that control the assembly of instructions into object files. these directives give the ability to include arbitrary data, align data, export symbols, switch sections, define constants and emit metadata.

the following table lists glyph assembler directives:

Directive	Arguments	Description
Data directives		
<code>.byte</code>	<i>expression-list</i>	8-bit comma separated words
<code>.short</code>	<i>expression-list</i>	16-bit comma separated words
<code>.long</code>	<i>expression-list</i>	32-bit comma separated words
<code>.quad</code>	<i>expression-list</i>	64-bit comma separated words
<code>.octa</code>	<i>expression-list</i>	128-bit comma separated words
<code>.string</code>	<i>"string"</i>	emit string
<code>.zero</code>	<i>integer</i>	emit zeroes
Alignment directives		
<code>.align</code>	<i>pow2</i> [<i>pad_val=0</i>] [<i>max</i>]	align to power of 2
<code>.balign</code>	<i>bytes</i> [<i>pad_val=0</i>]	byte align
Symbol directives		
<code>.globl</code>	<i>symbol_name, const_name</i>	emit symbol (<i>global scope</i>)
<code>.local</code>	<i>symbol_name, const_name</i>	emit symbol (<i>local scope</i>)
Section directives		
<code>.text</code>		emit <code>.text</code> section or make current
<code>.const</code>		emit <code>.const</code> section or make current
<code>.data</code>		emit <code>.data</code> section or make current
<code>.rodata</code>		emit <code>.rodata</code> section or make current
<code>.bss</code>		emit <code>.bss</code> section or make current
<code>.common</code>	<i>symbol_name, size, align</i>	emit common object to <code>.bss</code> section
<code>.section</code>	<i>section_name</i>	emit section (default <code>.text</code>) or make current
Miscellaneous directives		
<code>.equ</code>	<i>name, value</i>	constant definition
<code>.file</code>	<i>"filename"</i>	emit filename symbol
<code>.ident</code>	<i>"string"</i>	emit identification string
<code>.size</code>	<i>symbol, symbol</i>	emit symbol size
<code>.type</code>	<i>symbol, @function</i>	emit symbol type

Table 3.1.: Assembler Directives

3.4. Pseudo-instructions

the assembler implements a number of convenience psuedo-instruction aliases that are formed from regular instructions, but have implicit or deduced arguments.

the following table lists glyph assembler pseudo instruction aliases:

Pseudo-instruction	Expansion	Description
<code>nop</code>	<code>or.i64 r0,r0,r0</code>	no-operation
<code>li rc, expression</code>	(several expansions)	load immediate
<code>la rc, symbol</code>	(several expansions)	load address
<code>call symbol</code>	<code>jalib ibcall-reloc(text-label,const-label)</code>	procedure call
<code>ret</code>	<code>jtlib ibret-reloc(block-entry-label)</code>	procedure return
<code>cmp.lt.i64 rc, rb</code>	<code>compare.i64 rc, rb, lt</code>	compare less than (signed)
<code>cmp.ge.i64 rc, rb</code>	<code>compare.i64 rc, rb, ge</code>	compare greater or equal (signed)
<code>cmp.eq.i64 rc, rb</code>	<code>compare.i64 rc, rb, eq</code>	compare equal (signed)
<code>cmp.ne.i64 rc, rb</code>	<code>compare.i64 rc, rb, ne</code>	compare not equal (signed)
<code>cmp.ltu.i64 rc, rb</code>	<code>compare.i64 rc, rb, ltu</code>	compare less than (unsigned)
<code>cmp.geu.i64 rc, rb</code>	<code>compare.i64 rc, rb, geu</code>	compare greater or equal (unsigned)
<code>cmov.i64 rc, rb</code>	<code>compare.i64 rc, rb, mov</code>	conditional move
<code>mov.i64 rc, rb</code>	<code>logic.i64 rc, rb, mov</code>	copy register
<code>not.i64 rc, rb</code>	<code>logic.i64 rc, rb, not</code>	logical not
<code>neg.i64 rc, rb</code>	<code>logic.i64 rc, rb, neg</code>	signed negate
<code>bswap.i64 rc, rb</code>	<code>logic.i64 rc, rb, bswap</code>	byte swap
<code>ctz.i64 rc, rb</code>	<code>logic.i64 rc, rb, ctz</code>	count trailing zeros
<code>clz.i64 rc, rb</code>	<code>logic.i64 rc, rb, clz</code>	count leading zeroes
<code>ctpop.i64 rc, rb</code>	<code>logic.i64 rc, rb, ctpop</code>	count population

Table 3.2.: Assembler Directives

3.5. Calling convention

3.5.1. 16-bit calling convention

the 16-bit instruction packet, while intended to be used in conjunction with the 32-bit opcodes, is designed as a complete subset, so there is an ABI variant that targets a subset using only the 16-bit opcodes.

the register assignment for the 16-bit subset was chosen with this rationale:

- 2 blocks of 4 contiguous non-volatile *callee-save* and volatile *caller-save* registers.
- 3 special registers, 2 argument registers, 1 temporary register, and 3 save registers.
- 3 save registers to avoid excessive spilling around function calls.
- 1 temporary register to avoid spilling arguments to free a temporary.

the calling convention for the 16-bit subset is as follows:

- *immediate base* **ib** is set by **call** instructions and must point to a valid immediate block on function entry. function symbols are exported with two labels; one in the **.text** section, and one in the **.const** section. *immediate base* must be restored to the entry value in the function *epilogue* before it can be restored by **ret**.
- *argument registers* **a0** and **a1** are used for the first two arguments, and the remaining arguments are passed on the stack. *return value* is places in **a0** and **a1**, *temporary register* **t0** is a volatile register, and *frame pointer* (if enabled) uses **s0**. there are two more non-volatile *callee-save* registers, **s1** and **s2**.

the following table outlines the 16-bit register allocation showing register name alias, description, and non-volatile *callee-save* or volatile *caller-save* status.

name	alias	description	save
ib		immediate base	callee
r0	sp	stack pointer	callee
r1	s0/fp	saved register 0 / frame pointer	callee
r2	s1	saved register 1	callee
r3	s2	saved register 2	callee
r4	t0	temporary register 0	caller
r5	a0	argument register 0	caller
r6	a1	argument register 1	caller
r7	ra	return address / (pc,ib) link vector	caller

Table 3.3.: 16-bit register assignment

A. Appendix

A.1. 16-bit opcodes

15	13	12	10	9	7	6	2	1	0
uimm						00000	00		
simm						00001	00		
simm						00010	00		
simm						00011	00		
rc	uimm					00100	00		
rc	uimm					00101	00		
rc	uimm					00110	00		
rc	simm					00111	00		
rc	simm					01000	00		
rc	uimm					01001	00		
rc	uimm					01010	00		
rc	uimm					01011	00		
rc	uimm					01100	00		
rc	uimm					01101	00		
rc	uimm					01110	00		
rc	uimm					01111	00		
rc	rb	uimm				10000	00		
rc	rb	uimm				10001	00		
rc	rb	uimm				10010	00		
rc	rb	uimm				10011	00		
rc	rb	ra				10100	00		
rc	rb	ra				10101	00		
rc	rb	ra				10110	00		
rc	rb	ra				10111	00		
rc	rb	ra				11000	00		
rc	rb	ra				11001	00		
rc	rb	ra				11010	00		
rc	rb	ra				11011	00		
rc	rb	ra				11100	00		
rc	rb	ra				11101	00		
rc	rb	ra				11110	00		
uimm						11111	00		

break	uimm9
j	simm9
b	simm9
ibj	simm9
jalib.i64	rc, ib64(uimm6*8)
jtlib.i64	rc, ib64(uimm6*8)
movib.i64	rc, ib64(uimm6*8)
movi.i64	rc, simm6
addi.i64	rc, simm6
srli.i64	rc, uimm6
srai.i64	rc, uimm6
slli.i64	rc, uimm6
addib.i64	rc, ib32(uimm6*4)
leapc.i64	rc, ib32(uimm6*4)(pc)
loadpc.i64	rc, ib32(uimm6*4)(pc)
storepc.i64	rc, ib32(uimm6*4)(pc)
load.i64	rc, (uimm3*8)(rb)
store.i64	rc, (uimm3*8)(rb)
compare.i64	rc, rb, fun3
logic.i64	rc, rb, fun3
pin.i64	rc, rb, ra
and.i64	rc, rb, ra
or.i64	rc, rb, ra
xor.i64	rc, rb, ra
add.i64	rc, rb, ra
srl.i64	rc, rb, ra
sra.i64	rc, rb, ra
sll.i64	rc, rb, ra
sub.i64	rc, rb, ra
mul.i64	rc, rb, ra
div.i64	rc, rb, ra
illegal	uimm9

break uimm9
j simm9
b simm9
ibj simm9
jalib.i64 rc, ib64(uimm6*8)
jtlib.i64 rc, ib64(uimm6*8)
movib.i64 rc, ib64(uimm6*8)
movi.i64 rc, simm6
addi.i64 rc, simm6
srli.i64 rc, uimm6
srai.i64 rc, uimm6
slli.i64 rc, uimm6
addib.i64 rc, ib32(uimm6*4)
leapc.i64 rc, ib32(uimm6*4)(pc)
loadpc.i64 rc, ib32(uimm6*4)(pc)
storepc.i64 rc, ib32(uimm6*4)(pc)
load.i64 rc, (uimm3*8)(rb)
store.i64 rc, (uimm3*8)(rb)
compare.i64 rc, rb, fun3
logic.i64 rc, rb, fun3
pin.i64 rc, rb, ra
and.i64 rc, rb, ra
or.i64 rc, rb, ra
xor.i64 rc, rb, ra
add.i64 rc, rb, ra
srl.i64 rc, rb, ra
sra.i64 rc, rb, ra
sll.i64 rc, rb, ra
sub.i64 rc, rb, ra
mul.i64 rc, rb, ra
div.i64 rc, rb, ra
illegal uimm9